


Chapter 5

Machine Learning Basics

Learning Algorithms

- (Mitchell, 1997) A computer program is said to learn from *experience* E with respect to some class of *tasks* T and *performance measure* P , if its performance at tasks in T , as measured by P , improves with experience E
- Example: Linear Regression



- Task T : To predict y from x by outputting

$$\hat{y} = \mathbf{w}^T \phi(\mathbf{x}) = \phi(\mathbf{x})^T \mathbf{w}$$

- Experience E : To learn \mathbf{w} by minimizing, over a training set

$$(\mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})}),$$

$$\text{MSE}^{(\text{train})} = \frac{1}{m^{(\text{train})}} \|\hat{\mathbf{y}}^{(\text{train})} - \mathbf{y}^{(\text{train})}\|_2^2$$

where

$$\hat{\mathbf{y}}^{(\text{train})} = \Phi^{(\text{train})} \mathbf{w}, \quad \Phi^{(\text{train})} = \begin{bmatrix} \phi(\mathbf{x}_0^{(\text{train})})^T \\ \phi(\mathbf{x}_1^{(\text{train})})^T \\ \vdots \\ \phi(\mathbf{x}_{m-1}^{(\text{train})})^T \end{bmatrix}$$

$$\mathbf{y}^{(\text{train})} = (y_0^{(\text{train})}, y_1^{(\text{train})}, \dots, y_{m-1}^{(\text{train})})$$


- Performance P : To measure mean squared error on a test set $(\mathbf{X}^{(\text{test})}, \mathbf{y}^{(\text{test})})$, i.e.,

$$\text{MSE}^{(\text{test})} = \frac{1}{m^{(\text{test})}} \|\hat{\mathbf{y}}^{(\text{test})} - \mathbf{y}^{(\text{test})}\|_2^2$$

- To minimize $\text{MSE}_{\text{train}}$, w can be solved by setting

$$\nabla_w \text{MSE}^{(\text{train})} = 0$$

- A geometrical view is to solve $\hat{y}^{(\text{train})}$ as the projection of $y^{(\text{train})}$ onto the column space of $\Phi^{(\text{train})}$



- We then have

$$w = \left(\Phi^{(\text{train})T} \Phi^{(\text{train})} \right)^{-1} \Phi^{(\text{train})T} y^{\text{train}}$$

- The present model can be extended to include a bias term b

$$\hat{y} = \mathbf{w}^T \phi(\mathbf{x}) + b = \tilde{\mathbf{w}}^T \tilde{\phi}(\mathbf{x}),$$

with

$$\tilde{\mathbf{w}} = \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}, \quad \tilde{\phi}(\mathbf{x}) = \begin{bmatrix} \phi(\mathbf{x}) \\ 1 \end{bmatrix}$$

- Prediction with a polynomial of degree 2:

$$\hat{y} = w_2 x^2 + w_1 x^1 + b = \tilde{\mathbf{w}}^T \tilde{\phi}(x),$$

where

$$\tilde{\mathbf{w}} = \begin{bmatrix} w_2 \\ w_1 \\ b \end{bmatrix}, \quad \tilde{\phi}(x) = \begin{bmatrix} \phi_2(x) \\ \phi_1(x) \\ 1 \end{bmatrix} = \begin{bmatrix} x^2 \\ x^1 \\ 1 \end{bmatrix}$$

Generalization


- The model's ability to perform well on *new, previously unseen* inputs
- Generalization error is defined to be the *expected value of the error* on a *new input* and is typically estimated by measuring the performance on a *test set* collected separately from the *training set*
- Examples (\mathbf{x}, y) in the training and test sets are assumed to be drawn independently from the same distribution, $p_{\text{data}}(\mathbf{x}, y)$
- **Bayes error:** Minimum generalization error achieved by an oracle model having knowledge of $p_{\text{data}}(\mathbf{x}, y)$
- E.g.: if $y = \mathbf{w}^T \phi(\mathbf{x}) + \varepsilon$ and ε is Gaussian noise independent of \mathbf{x} , Bayes error = $\text{Var}(\varepsilon)$ with MSE measure

Training Error and Test Error

- (Pitfall) At first glance, the expected test error should be the same as the expected training error for a given model, because the data in these sets are drawn from the same distribution
- In practice, we sample the training set, use it to train the model, and then sample the test set to measure test error
- Generally, test error \geq training error


Underfitting vs. Overfitting

- Two objectives to achieve in designing a model
 1. Make training error small to avoid **underfitting**
 2. Make gap between training and test error small to avoid **overfitting**
- Trade-off can be made by altering the *model capacity*, which refers broadly to a model's ability to fit a wide variety of functions
- Example: Fitting a polynomial model to quadratic data



Capacity and Error

Typical relationship between capacity and error

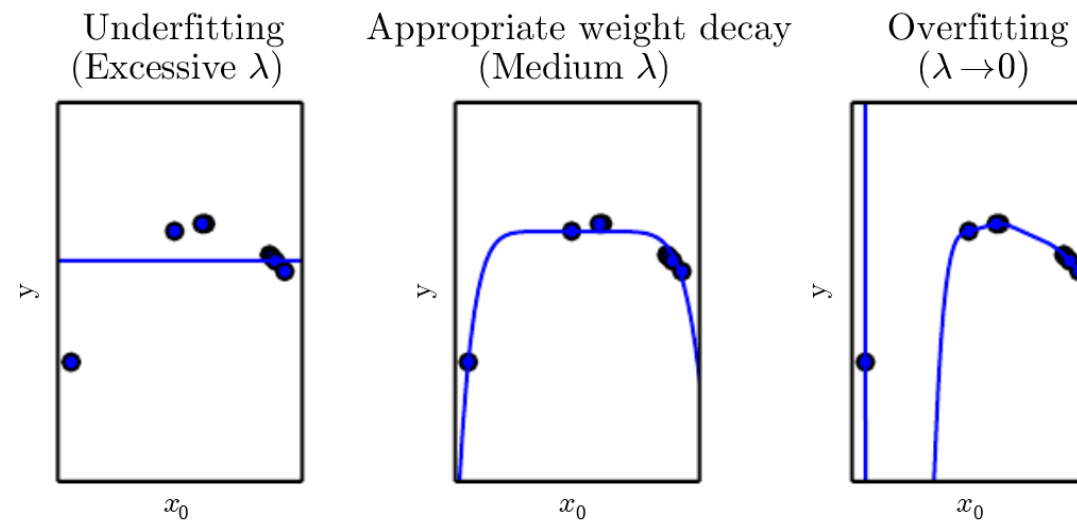


Regularization

- Modification made to the learning algorithm to reduce generalization error (usually at the cost of higher training error)
- Example: To include *weight decay* in the training criterion

$$J(\mathbf{w}) = \text{MSE}^{(\text{train})} + \lambda \mathbf{w}^T \mathbf{w}$$

where λ controls preference for small \mathbf{w} and is determined a priori



A degree-9 polynomial model fitted to quadratic data

Estimators

- **Point estimation:** To provide a single estimate of some quantity from observing independent and identically distributed (i.i.d.) samples
 - Consider m i.i.d samples $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ drawn from a Bernoulli distribution with mean θ

$$P(x^{(i)}; \theta) = \theta^{x^{(i)}} (1 - \theta)^{1-x^{(i)}}, \quad x^{(i)} = \{1, 0\}$$

- The *sample mean* can be used to give a point estimate of θ

$$\hat{\theta}_m = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

- A point estimator $\hat{\theta}_m$ of a parameter θ is any function of the observed samples $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$

$$\hat{\theta}_m = g(x^{(1)}, x^{(2)}, \dots, x^{(m)})$$

- θ is fixed but unknown from the **frequentist** viewpoint
- $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}$ are seen samples of a random variable
- As a result, $\hat{\theta}_m$ is a random variable

Bias

- The bias of the estimator $\hat{\theta}_m$ is defined as

$$\text{bias}(\hat{\theta}_m) = E(\hat{\theta}_m) - \theta,$$

where the expectation $E(\cdot)$ is taken w.r.t. $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}$

- $\hat{\theta}_m$ is *unbiased* if $\text{bias}(\hat{\theta}_m) = 0$
- $\hat{\theta}_m$ is *asymptotically unbiased* if $\lim_{m \rightarrow \infty} \text{bias}(\hat{\theta}_m) = 0$
- In the Bernoulli example, the sample mean is an unbiased estimator

$$E \left[\frac{1}{m} \sum_{i=1}^m x^{(i)} \right] = \frac{1}{m} \sum_{i=1}^m E \left[x^{(i)} \right] = \theta$$

Consistency

- An estimator $\hat{\theta}_m$ is said to be consistent *in probability* if

$$\lim_{m \rightarrow \infty} P\left(\left|\hat{\theta}_m - \theta\right| > \varepsilon\right) = 0, \quad \varepsilon > 0$$

- Consistency ensures that the bias of the estimator diminishes as the number of data samples grows
- In the Bernoulli example, the sample mean is consistent
 - Chebyshev's inequality

$$P(|X - \mu_X| > \varepsilon) \leq \frac{\sigma_x^2}{\varepsilon^2}$$

- $\hat{\theta}_m$ has mean θ , variance $\theta(1 - \theta)/m$

$$P(|\hat{\theta}_m - \theta| > \varepsilon) \leq \frac{\theta(1 - \theta)}{m\varepsilon^2} \Rightarrow \lim_{m \rightarrow \infty} P(|\hat{\theta}_m - \theta| > \varepsilon) = 0$$

Estimators for Gaussian Distribution

- Gaussian probability density function

$$\mathcal{N}(x^{(i)}; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2} \frac{(x^{(i)} - \mu)^2}{\sigma^2}\right)$$

- Sample mean (unbiased)

$$\hat{\mu}_m = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

- Sample variance (asymptotically unbiased)

$$\hat{\sigma}_m^2 = \frac{1}{m} \sum_{i=1}^m \left(x^{(i)} - \hat{\mu}_m\right)^2$$

- It can be shown that $E[\hat{\sigma}_m^2] = (m-1)\sigma^2/m$
- Unbiased sample variance $\tilde{\sigma}_m^2 = m\hat{\sigma}_m^2/(m-1)$

Variance of the Estimator

- Variance of the estimator indicates how much the estimator varies as a function of the samples; and its squared root is called standard error
- Example: Standard error of the sample mean $\hat{\mu}_m$

$$\text{SE}(\hat{\mu}_m) = \sqrt{\text{Var} \left[\frac{1}{m} \sum_{i=1}^m x^{(i)} \right]} = \frac{\sigma}{\sqrt{m}}$$

where σ is usually estimated by $\sqrt{\tilde{\sigma}_m^2}$

- By the central limit theorem,

$$\frac{\hat{\mu}_m - 0}{\text{SE}(\hat{\mu}_m)} \sim \mathcal{N}(0, 1)$$

- The 95 percent confidence interval can thus be derived as

$$(\hat{\mu}_m - 1.96\text{SE}(\hat{\mu}_m), \hat{\mu}_m + 1.96\text{SE}(\hat{\mu}_m))$$

- In experiments, it is common to say algorithm A performs better than B if its 95 percent upper bound of the test error is smaller than the lower bound of B's test error

Maximum Likelihood (ML) Estimation

- Consider examples $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}$ drawn independently from a distribution $p_{\text{model}}(\mathbf{x}; \boldsymbol{\theta})$, with parameter $\boldsymbol{\theta}$ being fixed but unknown
- The maximum likelihood estimator $\boldsymbol{\theta}_{\text{ML}}$ for $\boldsymbol{\theta}$ is defined as

$$\begin{aligned}\boldsymbol{\theta}_{\text{ML}} &= \arg \max_{\boldsymbol{\theta}} p_{\text{model}}(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(m)}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} \prod_{i=1}^m p_{\text{model}}(\mathbf{x}^{(i)}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^m \log p_{\text{model}}(\mathbf{x}^{(i)}; \boldsymbol{\theta})\end{aligned}$$

where $\sum_{i=1}^m \log p_{\text{model}}(\mathbf{x}^{(i)}; \boldsymbol{\theta})$ is the *log-likelihood function* of $\boldsymbol{\theta}$

- The sum over examples can be written as expectation w.r.t. the

empirical data distribution \hat{p}_{data}

$$\theta_{\text{ML}} = \arg \max_{\theta} E_{\mathbf{x} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}; \theta)$$

- Maximizing the log-likelihood amounts to minimizing the dissimilarity between the empirical data distribution \hat{p}_{data} (defined by the training set) and the model distribution p_{model} in terms of KL divergence:

$$D_{KL}(\hat{p}_{\text{data}} \| p_{\text{model}}) = E_{\mathbf{x} \sim \hat{p}_{\text{data}}} [\log \hat{p}_{\text{data}}(\mathbf{x}) - \log p_{\text{model}}(\mathbf{x})]$$

- Remarks
 - In information theory, $D_{KL}(\hat{p}_{\text{data}} \| p_{\text{model}})$ denotes the extra amount of information (in bits when using \log_2 base) needed to send a message \mathbf{x} drawn from \hat{p}_{data} with a code optimized for messages drawn from p_{model}
 - The *cross-entropy* $H(\hat{p}_{\text{data}}, p_{\text{model}})$ between \hat{p}_{data} and p_{model} is defined as

$$-E_{\mathbf{x} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}; \theta)$$

- It is easy to show that

$$-E_{\mathbf{x} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}; \boldsymbol{\theta}) = H(\hat{p}_{\text{data}}) + D_{KL}(\hat{p}_{\text{data}} \| p_{\text{model}})$$

- Minimizing the cross-entropy $H(\hat{p}_{\text{data}}, p_{\text{model}})$ is equivalent to maximizing the log-likelihood function $E_{\mathbf{x} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{x}; \boldsymbol{\theta})$ and thus minimizing the $D_{KL}(\hat{p}_{\text{data}} \| p_{\text{model}})$

Conditional Log-Likelihood Estimation

- The ML estimator generalized to learn a conditional probability $p_{\text{model}}(\mathbf{y}|\mathbf{x}; \boldsymbol{\theta})$ in order to predict \mathbf{y} given \mathbf{x}

$$\begin{aligned}\boldsymbol{\theta}_{\text{ML}} &= \arg \max_{\boldsymbol{\theta}} p_{\text{model}}(\mathbf{Y}|\mathbf{X}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} \sum_{i=1}^m \log p_{\text{model}}(\mathbf{y}^{(i)}|\mathbf{x}^{(i)}; \boldsymbol{\theta}) \\ &= \arg \max_{\boldsymbol{\theta}} E_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{\text{data}}(\mathbf{x}, \mathbf{y})} \log p_{\text{model}}(\mathbf{y}|\mathbf{x}; \boldsymbol{\theta})\end{aligned}$$

- Example: Linear Regression

$$y = \hat{y}(\mathbf{x}; \mathbf{w}) + \varepsilon = \mathbf{w}^T \phi(\mathbf{x}) + \varepsilon$$

where $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ is independent of \mathbf{x}

- It can be shown that $p(y|\mathbf{x}; \mathbf{w}) = \mathcal{N}(y; \hat{y}(\mathbf{x}; \mathbf{w}), \sigma^2)$
- The conditional log-likelihood of $p_{\text{model}}(\mathbf{Y}|\mathbf{X}; \mathbf{w})$ is given by

$$\sum_{i=1}^m \log p(y^{(i)} | \mathbf{x}^i; \mathbf{w}) = -m \log \sigma - \frac{m}{2} \log(2\pi) - \sum_{i=1}^m \frac{\|\hat{\mathbf{y}}^i - \mathbf{y}^{(i)}\|^2}{2\sigma^2}$$

- Maximizing the log-likelihood w.r.t. \mathbf{w} leads to the same problem of minimizing

$$\text{MSE}^{(\text{train})} = \frac{1}{m^{(\text{train})}} \|\hat{\mathbf{y}}^{(\text{train})} - \mathbf{y}^{(\text{train})}\|_2^2$$

- Therefore, the data model $y = \hat{y}(\mathbf{x}; \mathbf{w}) + \varepsilon$ provides an alternative view of the linear regression problem

Bayesian Statistics

- Assume the true parameter θ is a random variable governed by a prior probability distribution $p(\theta)$
- The goal is to infer the posterior distribution $p(\theta|\mathbf{X})$ of θ by combining the prior $p(\theta)$ and the data likelihood $p(\mathbf{X}|\theta)$ via Bayes' rule:

$$p(\theta|\mathbf{X}) = \frac{p(\mathbf{X}|\theta)p(\theta)}{p(\mathbf{X})}$$

- Example: Linear Regression

$$y = \hat{y}(\mathbf{x}; \mathbf{w}) + \varepsilon = \mathbf{w}^T \phi(\mathbf{x}) + \varepsilon$$

- Again, $p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(y; \hat{y}(\mathbf{x}; \mathbf{w}), \sigma^2)$, where presently $\sigma^2 = 1$
- Assume $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}_0, \boldsymbol{\Lambda}_0)$
- Consider \mathbf{x} to be deterministic data

- The posterior distribution $p(\mathbf{w}|\mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})})$ is given by

$$\begin{aligned}
 & p(\mathbf{w}|\mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})}) \\
 & \propto p(\mathbf{w})p(\mathbf{y}^{(\text{train})}|\mathbf{X}^{(\text{train})}, \mathbf{w}) \\
 & \propto \exp\left(-\frac{1}{2}(\mathbf{w} - \mathbf{u}_0)^T \mathbf{\Lambda}_0^{-1}(\mathbf{w} - \mathbf{u}_0)\right) \exp\left(-\frac{1}{2}(\mathbf{y} - \mathbf{\Phi}\mathbf{w})^T(\mathbf{y} - \mathbf{\Phi}\mathbf{w})\right) \\
 & \propto \exp\left(-\frac{1}{2}(\mathbf{w} - \mathbf{u}_m)^T \mathbf{\Lambda}_m^{-1}(\mathbf{w} - \mathbf{u}_m)\right) \\
 & = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}_m, \mathbf{\Lambda}_m)
 \end{aligned}$$

where

$$\begin{aligned}
 \mathbf{\Phi} &= \mathbf{\Phi}^{(\text{train})} \\
 \mathbf{\Lambda}_m &= (\mathbf{\Phi}^T \mathbf{\Phi} + \mathbf{\Lambda}_0^{-1})^{-1} \\
 \mathbf{u}_m &= \mathbf{\Lambda}_m (\mathbf{\Phi}^T \mathbf{y}^{(\text{train})} + \mathbf{\Lambda}_0^{-1} \mathbf{u}_0)
 \end{aligned}$$

- When $\boldsymbol{\mu}_0 = \mathbf{0}$ and $\mathbf{\Lambda}_0 = \frac{1}{\lambda} \mathbf{I}$, $\boldsymbol{\mu}_m$ leads to the same solution as

minimizing

$$J(\mathbf{w}) = \text{MSE}^{(\text{train})} + \lambda \mathbf{w}^T \mathbf{w}$$

- Given $p(\mathbf{w} | \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})})$, one can infer the probability distribution of $y^{(\text{new})}$ at unseen $\mathbf{x}^{(\text{new})}$ by

$$\begin{aligned} & p(y^{(\text{new})} | \mathbf{x}^{(\text{new})}, \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})}) \\ &= \int p(y^{(\text{new})}, \mathbf{w} | \mathbf{x}^{(\text{new})}, \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})}) d\mathbf{w} \\ &= \int p(\mathbf{w} | \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})}) p(y^{(\text{new})} | \mathbf{x}^{(\text{new})}, \mathbf{w}) d\mathbf{w} \\ &= \int \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}_m, \boldsymbol{\Lambda}_m) \mathcal{N}(y^{(\text{new})}; \hat{y}(\mathbf{x}^{(\text{new})}; \mathbf{w}), 1) d\mathbf{w} \end{aligned}$$

- Equivalently, this is to ask about the distribution of

$$y^{(\text{new})} = \hat{y}(\mathbf{x}^{(\text{new})}; \mathbf{w}) + \varepsilon = \mathbf{w}^T \phi(\mathbf{x}^{(\text{new})}) + \varepsilon$$

given $(\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})})$, with

$$p(\mathbf{w}|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})}) = \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}_m, \boldsymbol{\Lambda}_m)$$

$$p(\varepsilon|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})}) = \mathcal{N}(\varepsilon; 0, 1)$$

and \mathbf{w} , ε being conditionally independent

$$\begin{aligned} p(\mathbf{w}, \varepsilon|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})}) \\ = p(\mathbf{w}|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})})p(\varepsilon|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})}) \end{aligned}$$

- We further recognize that

1. $\mathbf{w}^T \phi(\mathbf{x}^{(\text{new})})|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})}$ is a Gaussian

$$\begin{aligned} \mathbf{w}^T \phi(\mathbf{x}^{(\text{new})})|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})} \\ \sim \mathcal{N}(\boldsymbol{\mu}_m^T \phi(\mathbf{x}^{(\text{new})}), \phi(\mathbf{x}^{(\text{new})})^T \boldsymbol{\Lambda}_m \phi(\mathbf{x}^{(\text{new})})) \end{aligned}$$

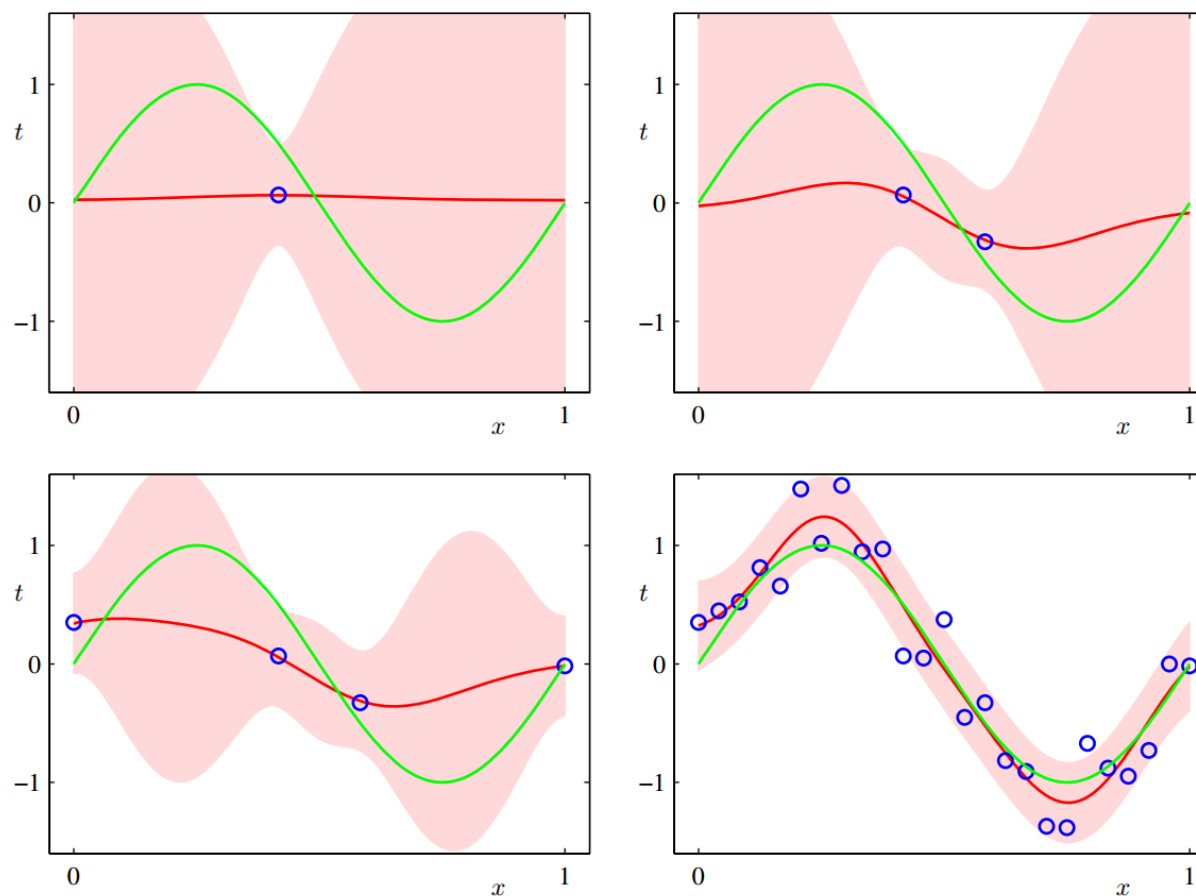
2. $\mathbf{w}^T \phi(\mathbf{x}^{(\text{new})}) + \varepsilon|\mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})}$ (sum of two conditionally

independent Gaussian's) is another Gaussian

$$\begin{aligned} & \mathbf{w}^T \phi(\mathbf{x}^{(\text{new})}) + \varepsilon | \mathbf{X}^{(\text{train})}, \mathbf{Y}^{(\text{train})} \\ & \sim \mathcal{N}(\boldsymbol{\mu}_m^T \phi(\mathbf{x}^{(\text{new})}), \phi(\mathbf{x}^{(\text{new})})^T \boldsymbol{\Lambda}_m \phi(\mathbf{x}^{(\text{new})}) + 1) \end{aligned}$$

- This concludes our evaluation for $p(y^{(\text{new})} | \mathbf{x}^{(\text{new})}, \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})})$

Bayesian Statistics



Ground truth (Green); $\mu_{y(\text{new})}$ (Red); Data (Blue); $\sigma_{y(\text{new})}$ (Pink)

Maximum A Posteriori (MAP) Estimation

- The full Bayesian treatment may sometimes be intractable
- To offer a point estimate in the Bayesian framework, we usually choose

$$\begin{aligned}\theta_{\text{MAP}} &= \arg \max_{\theta} p(\theta | \mathbf{X}) \\ &= \arg \max_{\theta} \frac{p(\mathbf{X} | \theta) p(\theta)}{p(\mathbf{X})} \\ &= \arg \max_{\theta} \log p(\mathbf{X} | \theta) + \log p(\theta)\end{aligned}$$

- Many regularized estimation strategies (e.g. ML+weight decay) can be interpreted as making the MAP inference, if the regularization term (e.g. weight decay) admits an explanation of $\log p(\theta)$
- Example: Linear Regression

$$\mathbf{w}_{\text{MAP}} = \arg \max_{\mathbf{w}} p(\mathbf{w} | \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})})$$

$$\begin{aligned}
&= \arg \max_{\mathbf{w}} \mathcal{N}(\mathbf{w}; \boldsymbol{\mu}_m, \boldsymbol{\Lambda}_m) \\
&= \boldsymbol{\mu}_m
\end{aligned}$$

- We may then choose the prediction model to be

$$\hat{y}(\mathbf{x}^{(\text{new})}; \mathbf{w}) = \mathbf{w}_{\text{MAP}}^T \phi(\mathbf{x}^{(\text{new})}) = \boldsymbol{\mu}_m^T \phi(\mathbf{x}^{(\text{new})}),$$

which in the present case coincides with the mean of the posterior distribution $p(y^{(\text{new})} | \mathbf{x}^{(\text{new})}, \mathbf{X}^{(\text{train})}, \mathbf{y}^{(\text{train})})$

- From the earlier derivation and definitions,

$$\boldsymbol{\mu}_m = \boldsymbol{\Lambda}_m (\boldsymbol{\Phi}^T \mathbf{y}^{(\text{train})} + \boldsymbol{\Lambda}_0^{-1} \mathbf{u}_0)$$

$$\boldsymbol{\Phi} = \begin{bmatrix} \phi(\mathbf{x}_0^{(\text{train})})^T \\ \phi(\mathbf{x}_1^{(\text{train})})^T \\ \vdots \\ \phi(\mathbf{x}_{m-1}^{(\text{train})})^T \end{bmatrix}, \quad \mathbf{y}^{(\text{train})} = \begin{bmatrix} y_0^{(\text{train})} \\ y_1^{(\text{train})} \\ \vdots \\ y_{m-1}^{(\text{train})} \end{bmatrix}$$

- Assuming $\mu_0 = 0$, we have

$$\begin{aligned}
 \hat{y}(\mathbf{x}^{(\text{new})}; \mathbf{w}) &= \phi(\mathbf{x}^{(\text{new})})^T \mathbf{\Lambda}_m \Phi^T \mathbf{y}^{(\text{train})} \\
 &= \sum_{i=0}^{m-1} \underbrace{\phi(\mathbf{x}^{(\text{new})})^T \mathbf{\Lambda}_m \phi(\mathbf{x}_i^{(\text{train})})}_{\text{Kernel}} y_i^{(\text{train})} \\
 &= \sum_{i=0}^{m-1} \underbrace{k(\mathbf{x}^{(\text{new})}, \mathbf{x}_i^{(\text{train})})}_{\text{Kernel}} y_i^{(\text{train})}
 \end{aligned}$$

where


$$k(\mathbf{x}^{(\text{new})}, \mathbf{x}_i^{(\text{train})}) = \phi(\mathbf{x}^{(\text{new})})^T \mathbf{\Lambda}_m \phi(\mathbf{x}_i^{(\text{train})})$$

- It is seen that the prediction $\hat{y}(\mathbf{x}^{(\text{new})}; \mathbf{w})$ is a weighted combination of the training data $y_i^{(\text{train})}$ with weights determined by the kernel function $k(\cdot)$ measuring a certain type of distance between $\mathbf{x}^{(\text{new})}$ and $\mathbf{x}_i^{(\text{train})}$


Support Vector Machines (SVM)

- One of the most influential approaches to supervised learning
- **Idea:** To find a hyperplane (in feature space) for classifying linearly separable training data according to the sign of $\hat{y}(\phi(\mathbf{x}))$

$$\hat{y}(\phi(\mathbf{x})) = \mathbf{w}^T \phi(\mathbf{x}) + b$$



raw data domain



feature domain

- The hyperplane, known as the decision boundary, is defined by

$$\hat{y}(\phi(\mathbf{x})) = \mathbf{w}^T \phi(\mathbf{x}) + b = 0$$

- \mathbf{w} is a vector orthogonal to every vector in the decision boundary
- Any point $\phi(\mathbf{x})$ to the decision boundary has a distance

$$\frac{|\hat{y}(\phi(\mathbf{x}))|}{\|\mathbf{w}\|}$$

- **Maximum margin classifiers:** Maximizing the smallest distance between the decision boundary and any of the training samples $\phi(\mathbf{x}_n)$

$$\arg \max_{\mathbf{w}, b} \min_n \frac{t_n \hat{y}(\phi(\mathbf{x}_n))}{\|\mathbf{w}\|} = \arg \max_{\mathbf{w}, b} \left\{ \frac{1}{\|\mathbf{w}\|} \min_n [t_n (\mathbf{w}^T \phi(\mathbf{x}_n) + b)] \right\}$$

where $t_n \in \{-1, 1\}$ is the ground-truth label associated with $\phi(\mathbf{x}_n)$

- Noting that \mathbf{w}, b can be scaled simultaneously ($\mathbf{w} \rightarrow \kappa \mathbf{w}, b \rightarrow \kappa b$) without changing the resulting distance, we can choose a κ such that

$t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) = 1$ for the closest point $\phi(\mathbf{x}_n)$ to the boundary

- This allows us to reformulate the problem as a *constrained optimization problem*

$$\min \frac{1}{2} \|\mathbf{w}\|_2^2, \text{ subject to}$$

$$t_n(\mathbf{w}^T \phi(\mathbf{x}_n) + b) \geq 1, \forall n$$

- Using the method of Lagrange multipliers, which will be introduced next, the solution for $\hat{y}(\phi(\mathbf{x}))$ can be solved as

$$\hat{y}(\phi(\mathbf{x})) = \sum_{i=1}^N a_n t_n \underbrace{\phi(\mathbf{x}_n)^T \phi(\mathbf{x})}_{\text{Kernel}} + b = \sum_{i=1}^N a_n t_n \underbrace{k(\mathbf{x}_n, \mathbf{x})}_{\text{Kernel}} + b$$

where $a_n \geq 0, \forall n$ and most of them are zero

- *Support vectors* refer to those $\phi(\mathbf{x}_n)$'s whose $a_n \neq 0$

Constrained Optimization

- To find the maximal/minimal value of $f(\mathbf{x})$ for \mathbf{x} (known as feasible points) in some set \mathbb{S} , e.g.

$$\arg \min_{\mathbf{x}} f(\mathbf{x}), \text{ subject to}$$

$$g^{(i)}(\mathbf{x}) = 0, \quad i = 1, \dots, m$$

$$h^{(j)}(\mathbf{x}) \leq 0, \quad j = 1, \dots, n$$

where $\mathbb{S} = \{\mathbf{x} | \forall i, g^{(i)}(\mathbf{x}) = 0, \forall j, h^{(j)}(\mathbf{x}) \leq 0\}$ are defined by m *equality constraints* and n *inequality constraints*

- The **Karush-Kuhn-Tucker (KKT)** approach obtains a solution by solving the unconstrained optimization of the Lagrangian function:

$$\min_{\mathbf{x}} \max_{\boldsymbol{\lambda}} \max_{\boldsymbol{\alpha}, \boldsymbol{\alpha} \geq \mathbf{0}} L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha})$$

where

$$L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha}) = f(\mathbf{x}) + \sum_{i=1}^m \lambda_i g^{(i)}(\mathbf{x}) + \sum_{j=1}^n \alpha_j h^{(j)}(\mathbf{x})$$

and $\boldsymbol{\lambda}$ and $\boldsymbol{\alpha}$ are termed Lagrange multipliers

- The optimal solution satisfies (necessary conditions)
 1. $\nabla L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha}) = \mathbf{0}$
 2. All constraints on \mathbf{x} and Lagrange multipliers are met
 3. Complementary slackness: $\boldsymbol{\alpha} \odot \mathbf{h}(\mathbf{x}) = \mathbf{0}$, i.e. $\alpha_j = 0$ for $h^{(j)}(\mathbf{x}) < 0$ (inactive) and $\alpha_j \geq 0$ for $h^{(j)}(\mathbf{x}) = 0$ (active)
- It is easy to see that when any constraint is violated,

$$\max_{\boldsymbol{\lambda}} \max_{\boldsymbol{\alpha}, \boldsymbol{\alpha} \geq \mathbf{0}} L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha}) = \infty,$$

which excludes infeasible points from being considered

- On the other hand, when the constraints are all satisfied,

$$\max_{\boldsymbol{\lambda}} \max_{\boldsymbol{\alpha}, \boldsymbol{\alpha} \geq \mathbf{0}} L(\boldsymbol{x}, \boldsymbol{\lambda}, \boldsymbol{\alpha}) = f(\boldsymbol{x}),$$

which ensures the optimum within feasible points is unchanged

- Example: To solve \boldsymbol{w}, b

$$\min \frac{1}{2} \|\boldsymbol{w}\|_2^2, \text{ subject to}$$

$$t_n(\boldsymbol{w}^T \phi(\boldsymbol{x}_n) + b) \geq 1, \forall n$$

- We set to zero the gradient w.r.t.

$$L(\boldsymbol{w}, b, \boldsymbol{\alpha}) = \frac{1}{2} \|\boldsymbol{w}\|_2^2 + \sum_{j=1}^n \alpha_j (1 - t_n(\boldsymbol{w}^T \phi(\boldsymbol{x}_n) + b))$$

and arrive at

$$\nabla_{\mathbf{w}} L(\mathbf{w}, b, \boldsymbol{\alpha}) = \mathbf{0} \Rightarrow \mathbf{w} = \sum_{j=1}^n \alpha_j t_n \phi(\mathbf{x}_n)$$

$$\nabla_b L(\mathbf{w}, b, \boldsymbol{\alpha}) = 0 \Rightarrow \sum_{j=1}^n \alpha_j t_n = 0$$

- At this point, we already have the form of the optimal \mathbf{w}
- To solve for $\boldsymbol{\alpha}$, we can substitute this \mathbf{w} back to the Lagrangian


$$\max_{\boldsymbol{\alpha}} \sum_{j=1}^n \alpha_j - \frac{1}{2} \sum_{p=1}^m \sum_{q=1}^m \alpha_p \alpha_q t_p t_q \phi(\mathbf{x}_p)^T \phi(\mathbf{x}_q), \text{ s.t.}$$

$$\alpha_i \geq 0, \forall i \quad \text{and} \quad \sum_{j=1}^n \alpha_j t_n = 0$$

- How to solve b ? (To be continued)

Principle Component Analysis (PCA)

- An *unsupervised learning* algorithm that learns a data representation



- Input: $\mathbf{x} \in \mathbb{R}^n$
- Representation of input: $\mathbf{c} \in \mathbb{R}^l$
- Decoder: $g(\mathbf{c}) = \mathbf{D}\mathbf{c}$ with \mathbf{D} composed of orthonormal columns
- Cost: $\|\mathbf{x} - g(\mathbf{c})\|_2^2$
- Encoder has the form $f(\mathbf{x}) = \mathbf{D}^T \mathbf{x}$ when Cost is minimized
- Objective: Given training data $\mathbf{X}^{(\text{train})}$, we wish to find \mathbf{D}

$$\arg \min_{\mathbf{D}} \|\mathbf{X}^{(\text{train})} - \mathbf{X}^{(\text{train})} \mathbf{D} \mathbf{D}^T\|_F^2, \text{ s.t. } \mathbf{D}^T \mathbf{D} = \mathbf{I}$$

- Recall that

$$\mathbf{X}^{(\text{train})} = \begin{bmatrix} \mathbf{x}_0^{(\text{train})T} \\ \mathbf{x}_1^{(\text{train})T} \\ \vdots \\ \mathbf{x}_{m-1}^{(\text{train})T} \end{bmatrix}, \quad \|\mathbf{A}\|_F^2 = \text{Tr}\{\mathbf{A}\mathbf{A}^T\} = \sum_{i,j} A_{i,j}^2$$

- By simple algebra, we have

$$\begin{aligned} & \|\mathbf{X}^{(\text{train})} - \mathbf{X}^{(\text{train})}\mathbf{D}\mathbf{D}^T\|_F^2 \\ &= \text{Tr}(\mathbf{X}^{(\text{train})}\mathbf{X}^{(\text{train})T}) - \text{Tr}(\mathbf{X}^{(\text{train})}\mathbf{D}\mathbf{D}^T\mathbf{X}^{(\text{train})T}) \end{aligned}$$

where the first term has nothing to do with \mathbf{D}

- The initial problem then simplifies to

$$\arg \max_{\mathbf{D}} \text{Tr}(\mathbf{X}^{(\text{train})}\mathbf{D}\mathbf{D}^T\mathbf{X}^{(\text{train})T}), \quad \text{s.t. } \mathbf{D}^T\mathbf{D} = \mathbf{I}$$

- Observing that the Trace operator has the property

$$\text{Tr}(\mathbf{ABC}) = \text{Tr}(\mathbf{BCA}) = \text{Tr}(\mathbf{CAB})$$

(as long as all matrix multiplications are allowed), we arrive at

$$\arg \max_{\mathbf{D}} \text{Tr}(\mathbf{D}^T \mathbf{X}^{(\text{train})T} \mathbf{X}^{(\text{train})} \mathbf{D}), \text{ s.t. } \mathbf{D}^T \mathbf{D} = \mathbf{I}$$

- By a further application of Singular Value Decomposition to

$$\mathbf{X}_{m \times n}^{(\text{train})} = \mathbf{U}_{m \times m} \mathbf{\Sigma}_{m \times n} \mathbf{V}_{n \times n}^T,$$

- \mathbf{U} is the eigenvector matrix of $\mathbf{X}_{m \times n}^{(\text{train})} \mathbf{X}_{m \times n}^{(\text{train})T}$ and satisfies

$$\mathbf{U} \mathbf{U}^T = \mathbf{U}^T \mathbf{U} = \mathbf{I}_{m \times m}$$

- \mathbf{V} is the eigenvector matrix of $\mathbf{X}_{m \times n}^{(\text{train})T} \mathbf{X}_{m \times n}^{(\text{train})}$ and satisfies

$$\mathbf{V} \mathbf{V}^T = \mathbf{V}^T \mathbf{V} = \mathbf{I}_{n \times n}$$

– Σ is the singular matrix given by

$$\left[\begin{array}{cccc|c} \sigma_1 & 0 & \dots & 0 & \\ 0 & \sigma_2 & \dots & 0 & \\ \vdots & \vdots & \ddots & 0 & \\ 0 & 0 & 0 & \sigma_r & \\ \hline & & \mathbf{0} & & \end{array} \right]_{m \times n}$$

- We see that the columns of D have an obvious choice of the l columns in V which corresponds to the largest l singular values

$$\arg \max_D \text{Tr}(D^T V \Sigma^T \Sigma V^T D), \text{ s.t. } D^T D = I$$

where

$$\Sigma^T \Sigma = \begin{bmatrix} \sigma_1^2 & 0 & \dots & 0 & \vdots & 0 \\ 0 & \sigma_2^2 & \dots & 0 & \vdots & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & \sigma_r^2 & \vdots & 0 \\ \hline & & & & \mathbf{0} & \mathbf{0} \end{bmatrix}_{n \times n}$$

Gradient-based Learning

- Learning algorithms often seek to minimize/maximize some objective function w.r.t. the model parameter, e.g.

$$\arg \min_{\mathbf{w}} J(\mathbf{w}) = -E_{\mathbf{x}, y \sim \hat{p}_{\text{data}}} [p_{\text{model}}(y|\mathbf{x})]$$

- Very often, there is no closed-form solution
- Gradient-based learning algorithms are thus called for to update estimates of the solution via an iterative procedure

Steepest Descent

- To decrease $J(\mathbf{w})$ in the direction in which it decreases the fastest

$$\mathbf{w}^{(n+1)} = \mathbf{w}^{(n)} - \epsilon \nabla_{\mathbf{w}} J(\mathbf{w}^{(n)})$$

where ϵ controls the step size for each update

- The negative gradient $-\nabla_{\mathbf{w}} J(\mathbf{w}^{(n)})$ points to the direction in which $J(\mathbf{w})$ decreases the fastest at $\mathbf{w}^{(n)}$
 - To see this, we define the directional derivative at \mathbf{w}_0 to be

$$\frac{\partial}{\partial \alpha} J(\mathbf{w}_0 + \alpha \mathbf{u})$$

- It can then be evaluated as

$$\frac{\partial}{\partial \alpha} J(\mathbf{w}_0 + \alpha \mathbf{u}) = \mathbf{u}^T \nabla_{\mathbf{w}} J(\mathbf{w}_0)$$

using the Taylor-1 approximation at \mathbf{w}_0

$$J(\mathbf{w}) \approx J(\mathbf{w}_0) + (\mathbf{w} - \mathbf{w}_0)^T \nabla_{\mathbf{w}} J(\mathbf{w}_0)$$

- The unit vector \mathbf{u} that points in the direction $-\nabla_{\mathbf{w}} J(\mathbf{w}_0)$ yields a minimal directive among other unit vectors
- Instead of using a fixed ϵ , we can use line search to adapt its value to the curvature of $J(\mathbf{w})$ at \mathbf{w}_0 along $-\nabla_{\mathbf{w}} J(\mathbf{w}_0)$
- This relies on approximating $J(\mathbf{w})$ at \mathbf{w}_0 with Taylor-2 approximation

$$J(\mathbf{w}) \approx J(\mathbf{w}_0) + (\mathbf{w} - \mathbf{w}_0)^T \nabla_{\mathbf{w}} J(\mathbf{w}_0) + \frac{1}{2}(\mathbf{w} - \mathbf{w}_0)^T \mathbf{H}(\mathbf{w} - \mathbf{w}_0)$$

where \mathbf{H} is the Hessian matrix defined as


$$\mathbf{H}(J(\mathbf{w}_0))_{i,j} = \frac{\partial^2}{\partial w_i \partial w_j} J(\mathbf{w}_0)$$

- We wish to find an ϵ that maximizes

$$J(\mathbf{w}_0) - J(\mathbf{w}_0 - \epsilon \nabla_{\mathbf{w}} J(\mathbf{w}_0))$$

- The optimal ϵ is given by

$$\epsilon^* = \frac{\nabla_{\mathbf{w}} J(\mathbf{w}_0)^T \nabla_{\mathbf{w}} J(\mathbf{w}_0)}{\nabla_{\mathbf{w}} J(\mathbf{w}_0)^T \mathbf{H} \nabla_{\mathbf{w}} J(\mathbf{w}_0)}$$



Newton's Method

- To solve for the critical point that minimizes $J(\mathbf{w})$ approximated by the Taylor-2 expansion at \mathbf{w}_0 as the new estimate

$$\arg \min_{\mathbf{w}} J(\mathbf{w}_0) + (\mathbf{w} - \mathbf{w}_0)^T \nabla_{\mathbf{w}} J(\mathbf{w}_0) + \frac{1}{2} (\mathbf{w} - \mathbf{w}_0)^T \mathbf{H} (\mathbf{w} - \mathbf{w}_0)$$

- By setting the gradient w.r.t. \mathbf{w} to zero, we have

$$\mathbf{w}^* = \mathbf{w}_0 - \mathbf{H}(J(\mathbf{w}_0))^{-1} \nabla_{\mathbf{w}} J(\mathbf{w}_0)$$

- The iterative update of \mathbf{w} then becomes

$$\mathbf{w}^{(n+1)} = \mathbf{w}^{(n)} - \mathbf{H}(J(\mathbf{w}^{(n)}))^{-1} \nabla_{\mathbf{w}} J(\mathbf{w}^{(n)})$$

Example

- We wish to find $\mathbf{w}^* = \arg \min_{\mathbf{w}} J(\mathbf{w})$, with

$$J(\mathbf{w}) = \frac{1}{4}(\mathbf{w}^T \mathbf{M} \mathbf{w})^2, \quad \mathbf{M} = \begin{bmatrix} 6 & -4 \\ -4 & 6 \end{bmatrix}$$


- \mathbf{M} has the following pairs of eigenvalues and eigenvectors


$$\lambda_1 = 1 \rightarrow \mathbf{v}_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}, \quad \lambda_2 = 5 \rightarrow \mathbf{v}_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

- The gradient and the Hessian matrix are given by

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = (\mathbf{w}^T \mathbf{M} \mathbf{w}) \mathbf{M} \mathbf{w}$$

$$\mathbf{H}\{J(\mathbf{w})\} = 2\mathbf{M}\mathbf{w}(\mathbf{M}\mathbf{w})^T + \mathbf{w}^T \mathbf{M} \mathbf{w} \mathbf{M}$$





Review

- Overfitting vs. Underfitting
- Generalization
- Regularization
- Estimators (ML vs. MAP)
- Support Vector Machines (SVM)
- Constrained Optimization
- Principle Component Analysis (PCA)
- Gradient-based Learning