

Introduction to Reinforcement Learning

I-Chen Wu

- Sutton, R.S. and Barto, A.G., Reinforcement Learning: An Introduction, MIT Press, Cambridge, MA, 1998.
 - <http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
 - Bible in this area.
- David Silver, Online Course for Deep Reinforcement Learning.
 - <http://www.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>



David Silver:
(the leader of the AlphaGo team)

“DL+RL = AI”


Many Faces of Reinforcement Learning

- Computer Science
 - Machine Learning
- Engineering
 - Optimal Control
- Mathematics
 - Operations Research
- Economics
 - Bounded Rationality
- Psychology
 - Classical/Operant Conditioning
- Neuroscience
 - Reward System

Branches of Machine Learning

- **Supervised Learning (SL)**
 - learning from a training set of labeled examples provided by a knowledgeable external supervisor.
- **Unsupervised Learning (UL)**
 - typically about finding structure hidden in collections of unlabeled data.
- **Reinforcement Learning (RL)**
 - learning from interaction

[Silver]



What are different from others?

- Characteristics:

- No supervisor, only a **reward** signal
- Feedback is delayed, not instantaneous
- Time really matters
- Agent's actions affect the subsequent data

- UL vs. RL:

- RL is learning from interaction.
- RL does not rely on examples of correct behavior.
- RL is trying to maximize a reward signal, instead of trying to find hidden structure.



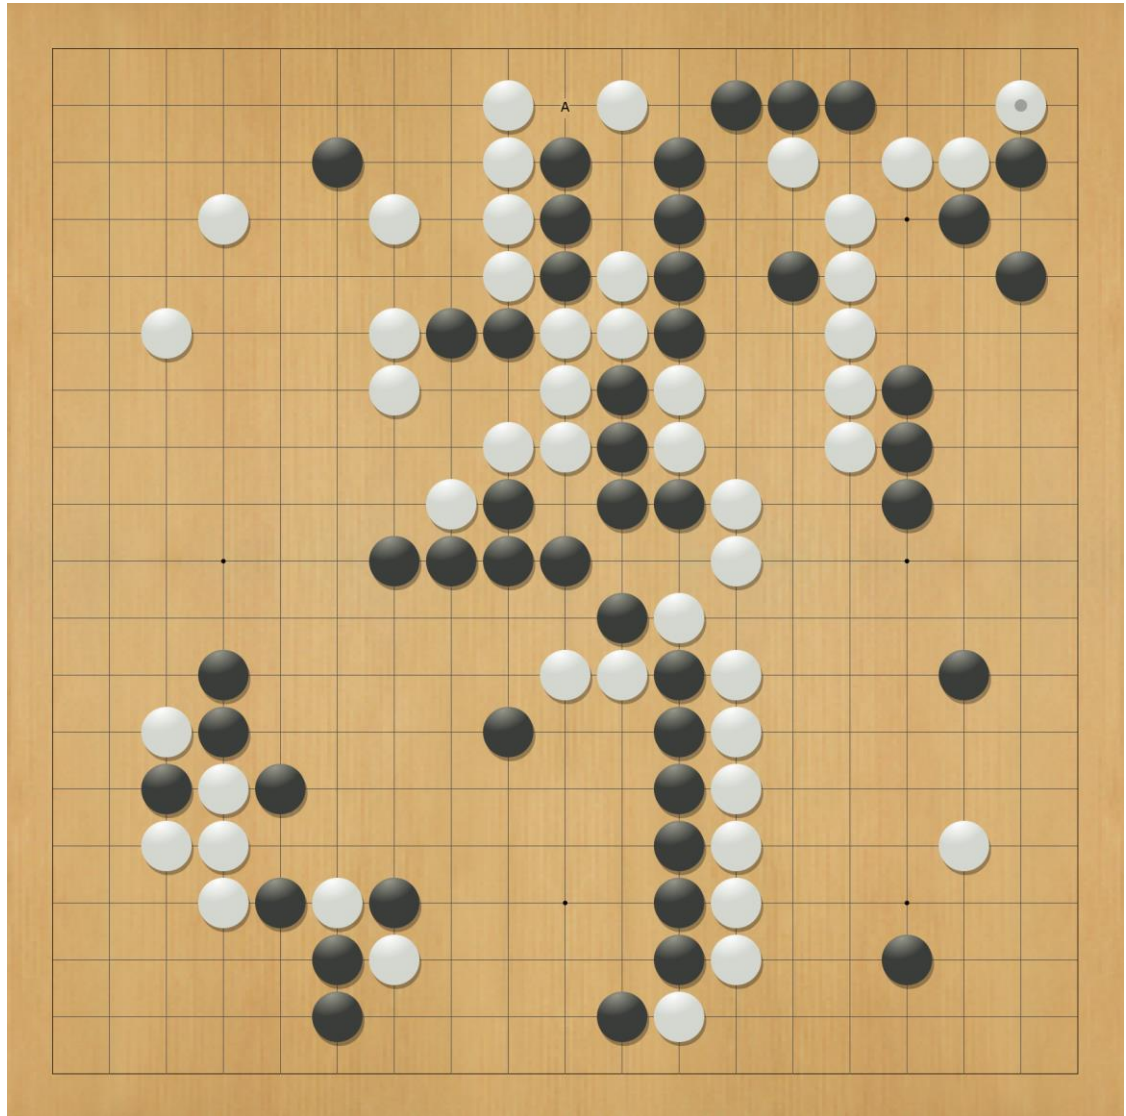
Successful Examples

- In AI, it has been used to defeat human champions at games of skill.
 - [Backgammon](#) (Tesauro, 1994).
 - [Connect6/2048/Threes!](#) (Wu et al., 2015). Reach the top levels.
 - [Go programs, used in the past 10 years. \(Monte-Carlo Tree Search\)](#)
 - [AlphaGo, using deep reinforcement learning \(2016\)](#)
 - In robotics, fly stunt maneuvers in [robot-controlled helicopters](#) (Abbeel et al.) and make a humanoid robot walk.
 - In economics, manage an investment portfolio (Choi et al.).
 - In neuroscience, [model the human brain](#) (Schultz et al.);
 - In psychology, [predict animal behavior](#) (Sutton and Barto).
 - In systems, control a power station
 - In engineering, it has been used to [allocate bandwidth to mobile phones](#) and to manage complex power systems (Ernst et al.).
- (Not even include successful examples for [deep reinforcement learning](#))



Board Game: Go

- Game 1: AlphaGo
vs. 李世石



Stochastic Game: 2048 (lab)

2	32768	8192	4096
16384	1024	512	256
2048	32	64	128
16	16	2	4

The First Game Reaching 65536 in the World (in 10,000 Trials)

<http://2048.aigames.nctu.edu.tw/replay.php>

2048

SCORE 1031392

BEST 1031392



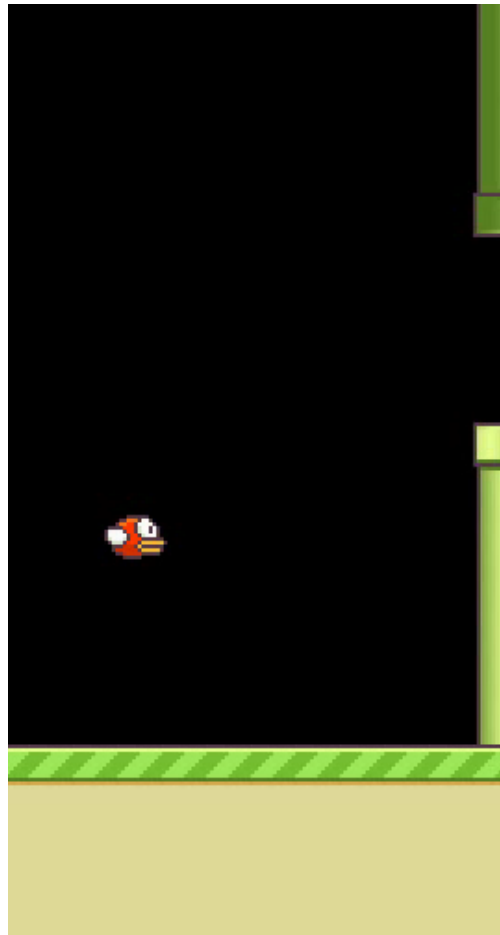
I-Chen Wu

Game over!

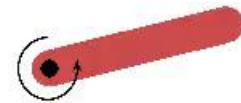
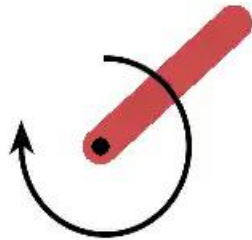
Try again

65536

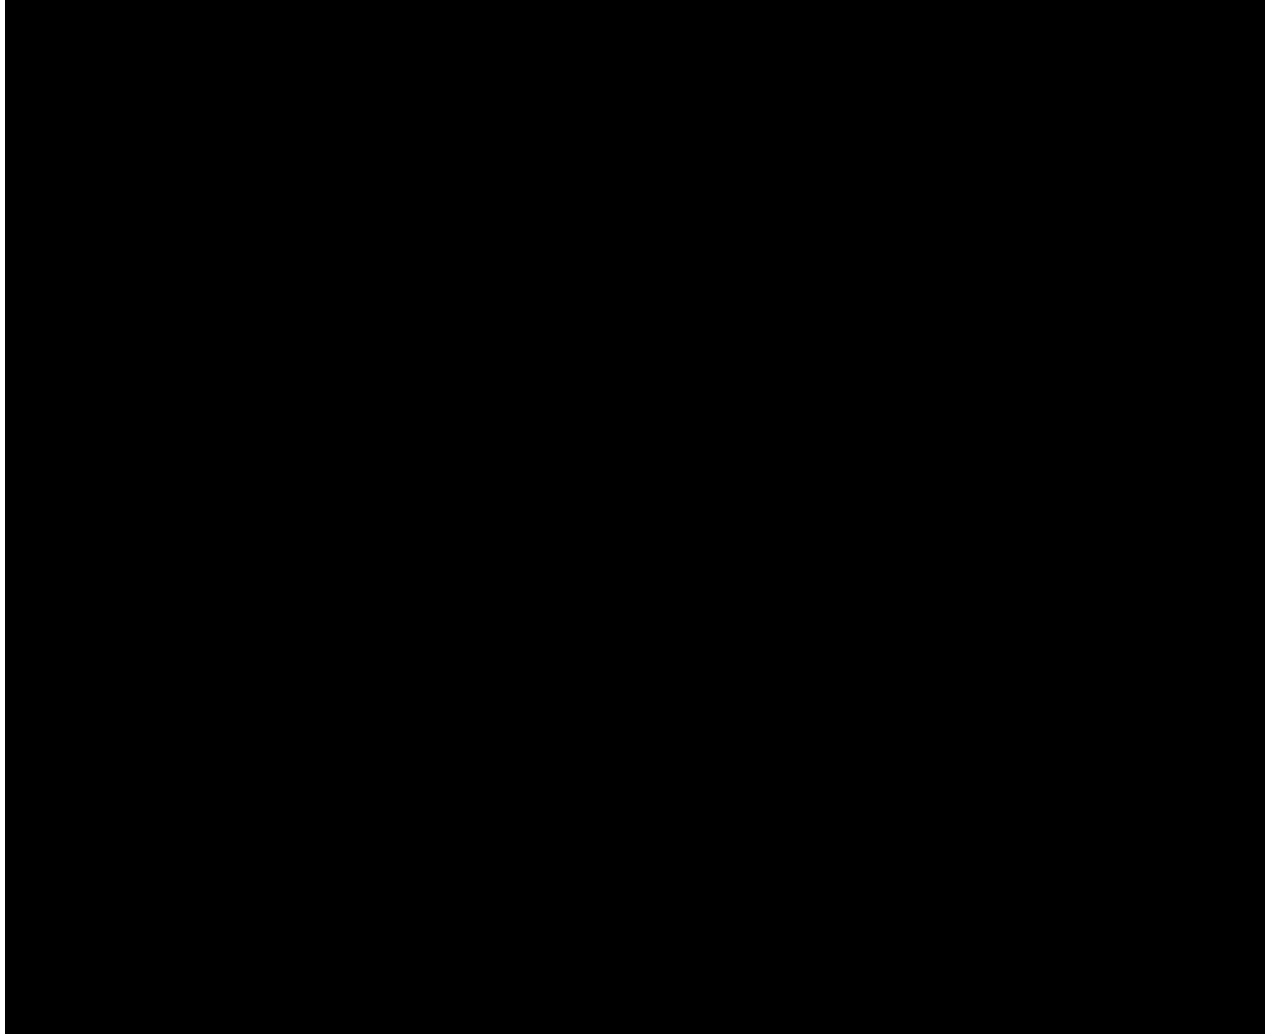
Video Games: Flappy Bird (lab)



Open AI: Pole Balancing (lab)



RL Demo



RL Demo

[Deisenroth et al, 2011] Learning to Control a Low-Cost Manipulator using Data-Efficient Reinforcement Learning

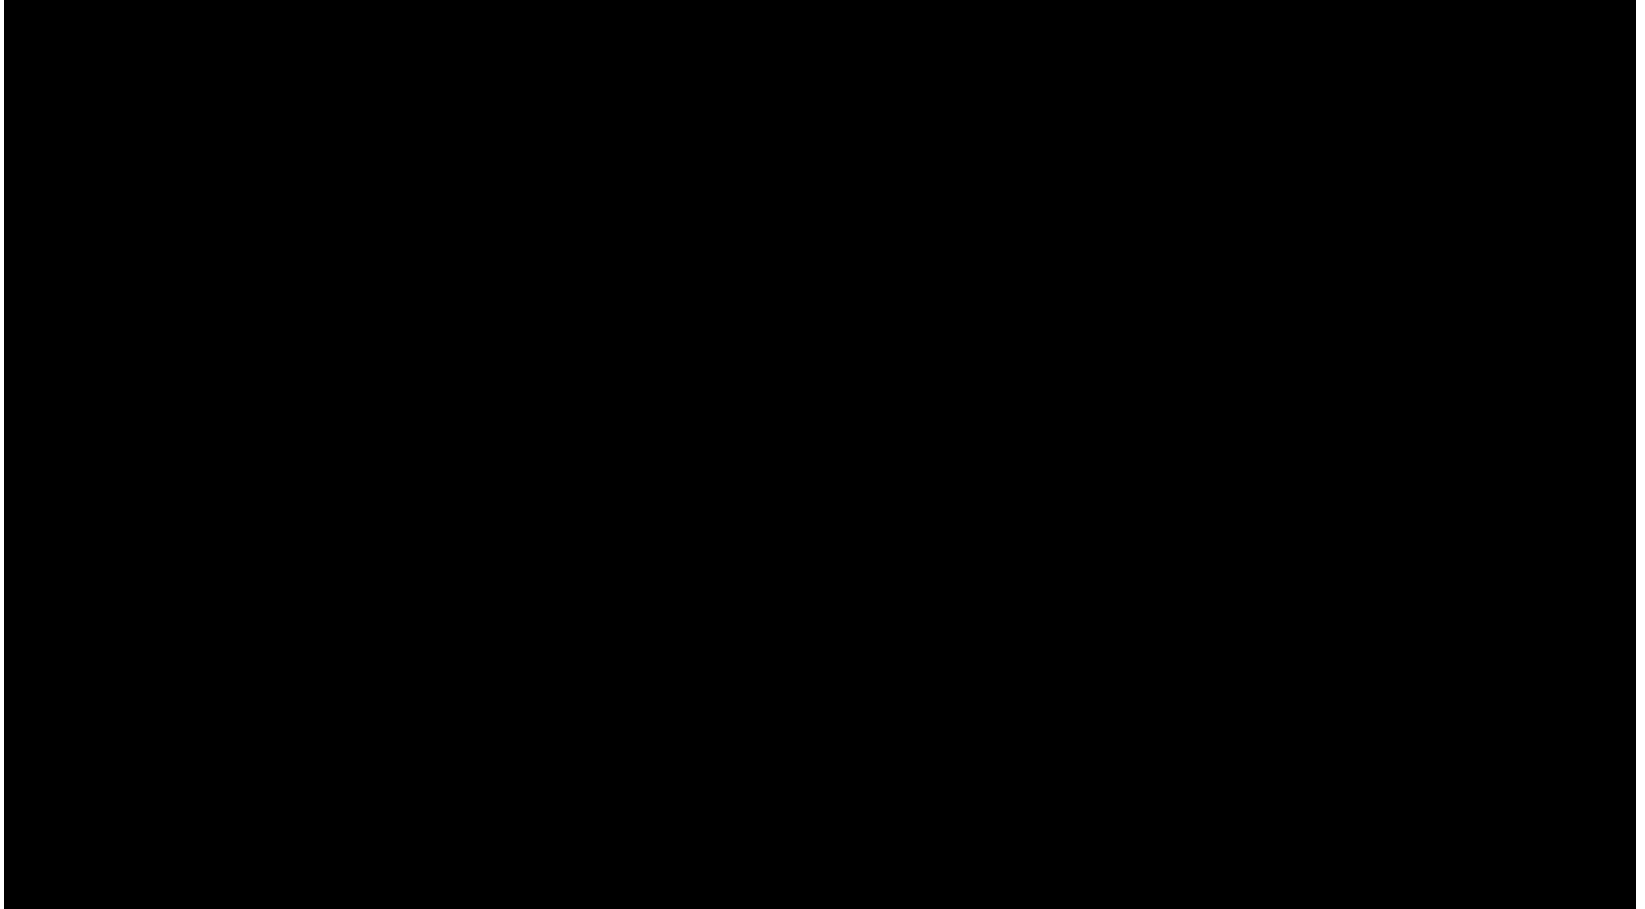
Marc Peter Deisenroth, Carl Edward Rasmussen, Dieter Fox

**Learning to Control a Low-Cost Robotic Manipulator
using Data-Efficient Reinforcement Learning**

R:SS 2011



Learning Contact-Rich Manipulation Skills with Guided Policy Search [Levine et. al. 2015]



ChatBot

- Hi, may I help you?
- I'm looking for a Chinese restaurant
 - Which area in mind?
- Somewhere in the downtown.
 - Hu Nan Restaurant is recommended by many people.
- Noop! The food is too hot.
 - How about Dumpling House?
- Good. Give me the direction to it.
 - It is located in
- Thank you.
 - Are you satisfied with the service.
- Yes.

Stock Market




Reinforcement Learning

- A **computational approach** to learning from **interaction**
 - Explore designs for machines that are effective in
 - ▶ solving learning problems of scientific or economic interest,
 - ▶ evaluating the designs through mathematical analysis or computational experiments.
 - Focus on **goal-directed learning from interaction**, when compared with other approaches to machine learning.
 - The learner must discover which actions yield the most reward by trying them.
 - ▶ Two characteristics: most important distinguishing features of reinforcement learning.
 - **trial-and-error search**
 - **delayed reward**


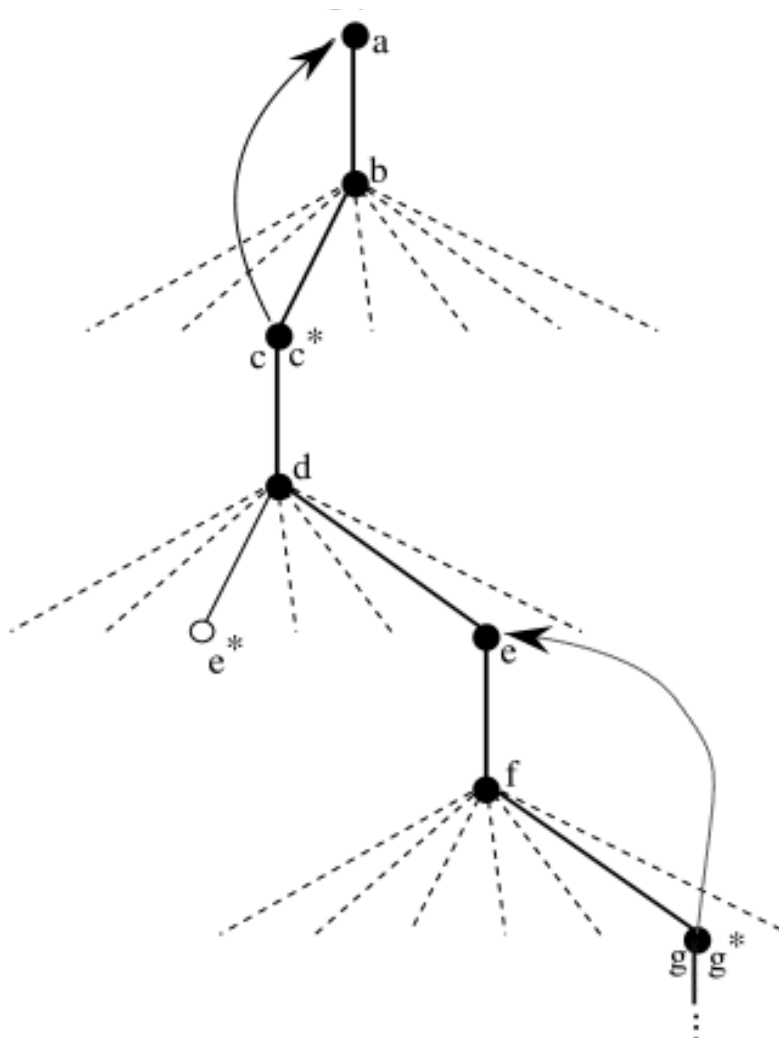



Agent-Environment Interaction Framework

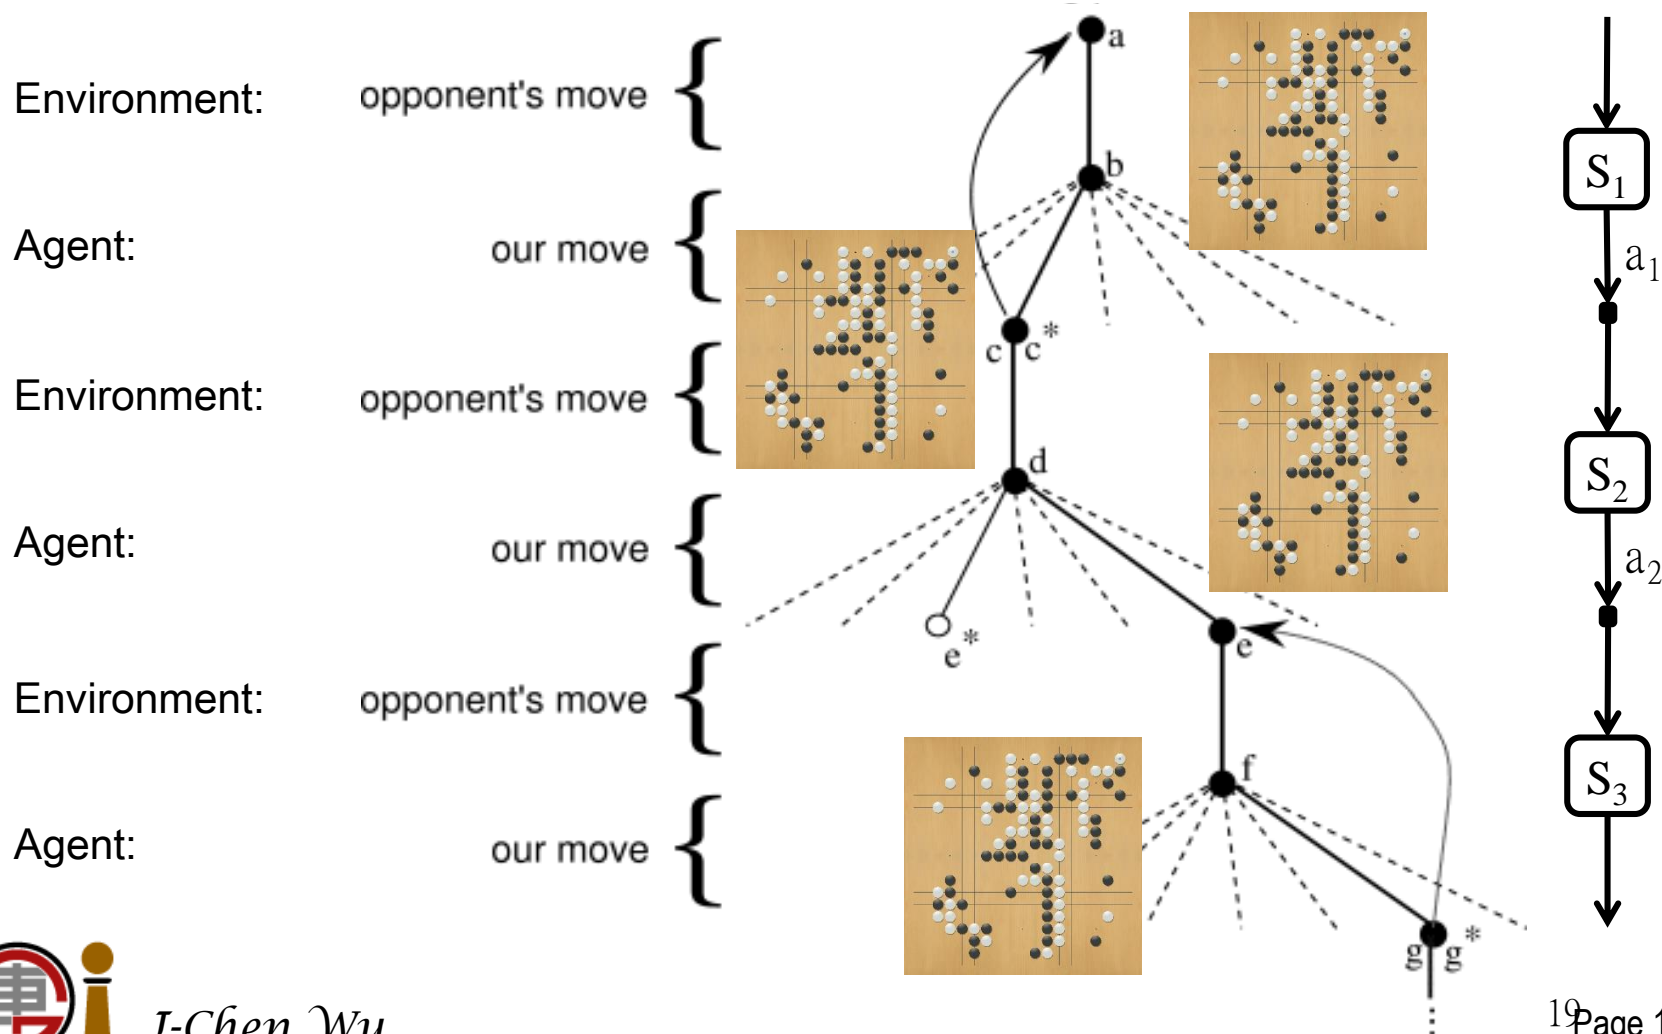
- **Agent**: The learner and decision-maker.
- **Environment**: The thing it interacts with, comprising everything outside the agent.
- **State**: whatever information is available to the agent.
- **Reward**: single numbers.



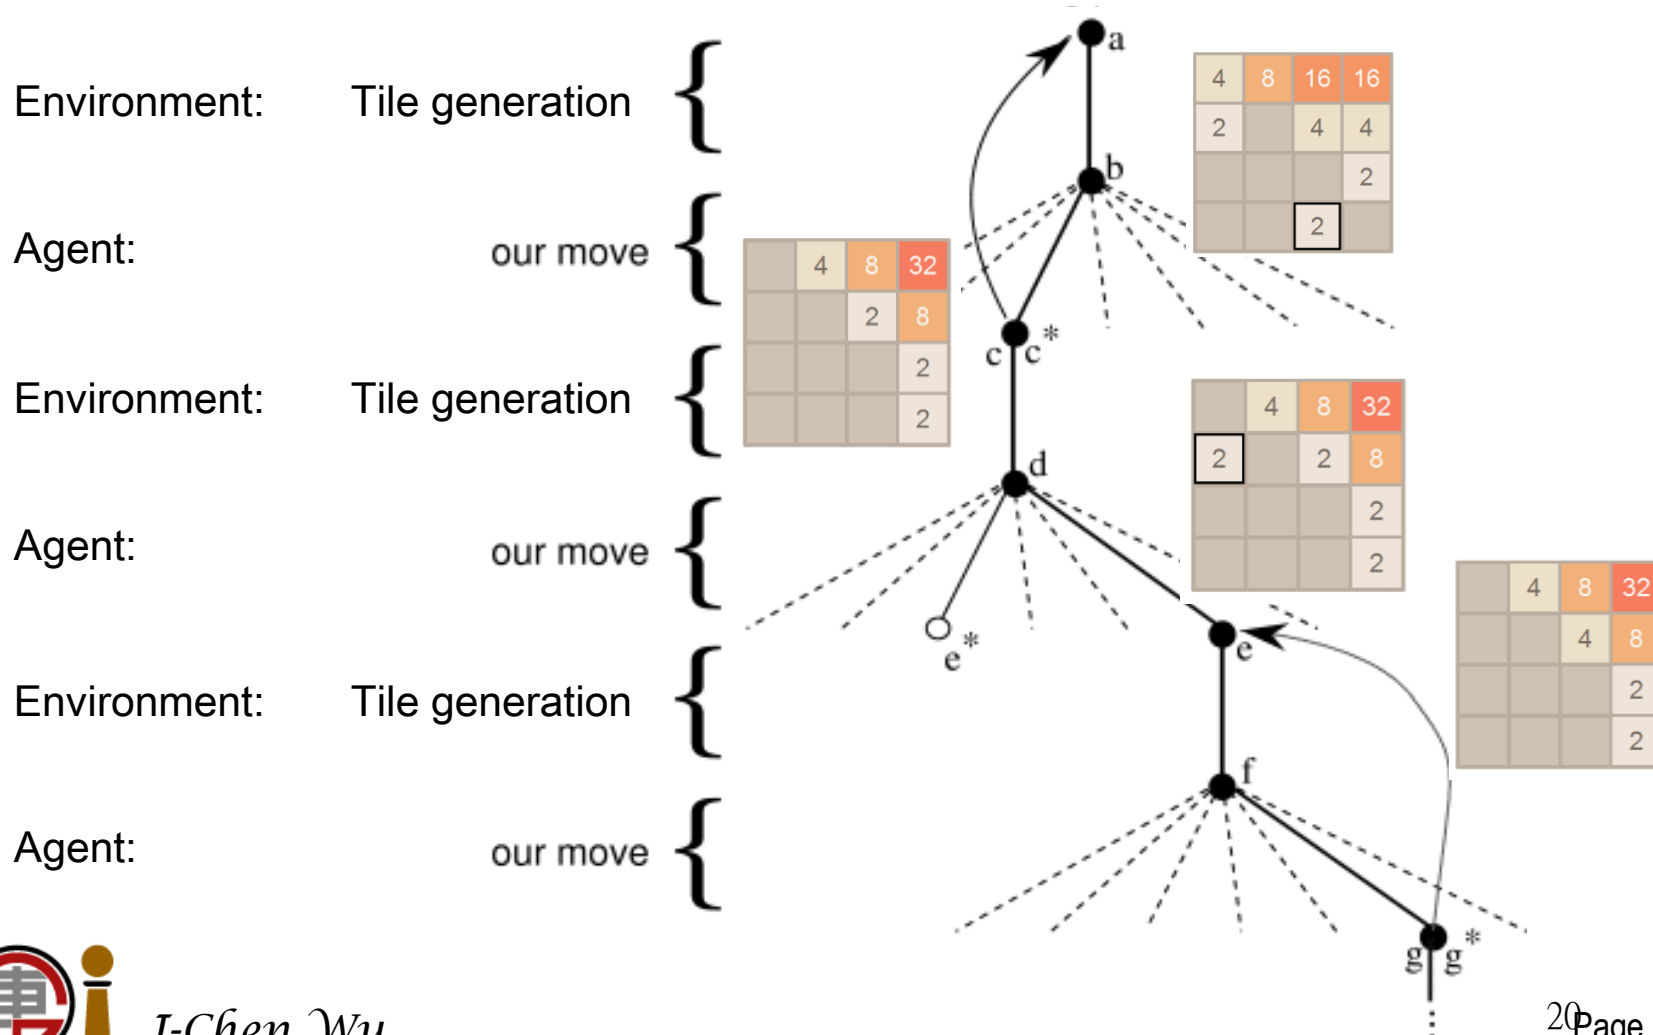
States and Actions in the Framework



Go



2048



Robot

Environment: Dynamics


Agent: Navigate

Environment: Dynamics

Agent: Navigate

Environment: Dynamics

Agent: Navigate




Markov Decision Processes (MDP)


- A **Markov Decision Process** is a tuple

$\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- \mathcal{S} is a finite set of states
- \mathcal{A} is a finite set of actions
- \mathcal{P} is a state transition probability matrix (part of the environment),
$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$
- \mathcal{R} is a reward function,
$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$
- γ is a discount factor $\gamma \in [0, 1]$.



Markov Property



- An **episode**: (assuming finite and MDP here for simplicity)

- States: S_i
 - ▶ Initial state: S_1
 - ▶ Current state: S_t
 - ▶ End state: S_T (not necessarily required)
- Actions: A_i
- **History**: $H_t = (S_1, A_1, R_2, S_2, A_2, R_3, S_3, \dots, R_t)$


- Markov Property:

- “The future is independent of the past given the present”
- A state S_t is **Markov** if and only if
$$\mathbb{P}[S_{t+1} | S_t] = \mathbb{P}[S_{t+1} | S_1, \dots, S_t]$$



Environment State vs. Agent State

- The **environment state S_t^e** :
 - the environment's private representation
 - ▶ i.e. whatever data the environment uses to pick the next observation/reward
 - The environment state is not necessarily visible to the agent
 - ▶ Even if S_t^e is visible, it may contain irrelevant information
- The **agent state S_t^a** :
 - The agent's internal representation
 - ▶ i.e. whatever information the agent uses to pick the next action
 - ▶ i.e. it is the information used by reinforcement learning algorithms
 - It can be any function of history:
$$S_t^a = f(H_t)$$
- **Partially Observable**: (not discussed here)
 - When $S_t^a \neq S_t^e$

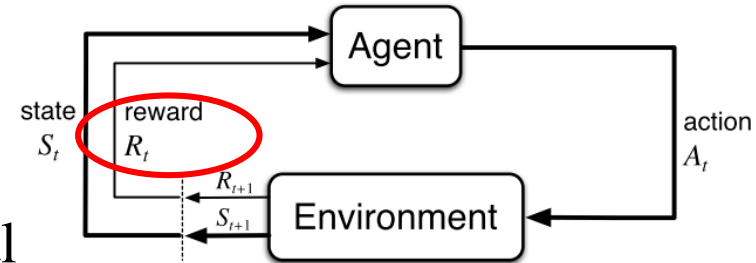


Example: Mahjong


- Partially observable:



Rewards



- A reward R_t is a **scalar feedback** signal
 - Indicates how well agent is doing at step t
 - The agent's job is to maximize cumulative reward
 - Reinforcement learning is based on the **reward hypothesis**
 - Example: (2048)



Definition (Reward Hypothesis)

- All goals can be described by the maximization of expected cumulative reward


Rewards for Previous Examples?

- In AI, it has been used to defeat human champions at games of skill.
 - Backgammon (Tesauro, 1994).
 - Connect6/2048/Threes! (Wu et al., 2015). Reach the top levels.
 - Go programs, used in the past 10 years. (Monte-Carlo Tree Search)
 - AlphaGo, using deep reinforcement learning (2016)
- In robotics, fly stunt maneuvers in robot-controlled helicopters (Abbeel et al.) and make a humanoid robot walk.
- In economics, manage an investment portfolio (Choi et al.).
- In neuroscience, model the human brain (Schultz et al.);
- In psychology, predict animal behavior (Sutton and Barto).
- In systems, control a power station
- In engineering, it has been used to allocate bandwidth to mobile phones and to manage complex power systems (Ernst et al.).



Sequential Decision Making

- Goal:
 - Select actions to maximize total future reward
- Maximize $R_{t+1} + R_{t+2} + \dots + R_T$
 - assuming time = t .



- Notes:
 - Actions may have long term consequences
 - Reward may be delayed
 - It may be better to sacrifice immediate reward to gain more long-term reward



Sequential Decision Making – Examples

- Examples:

- In 2048, establish a sequence of $(2^t, 2^{t-1}, 2^{t-2}, \dots)$
- In chess, block opponent moves to help winning chances many moves from now.
- In a financial investment, may take months to mature
- In robotics, refuel a helicopter to prevent a crash.


2	32768	8192	4096
16384	1024	512	256
2048	32	64	128
16	16	2	4

Major Components of an RL Agent

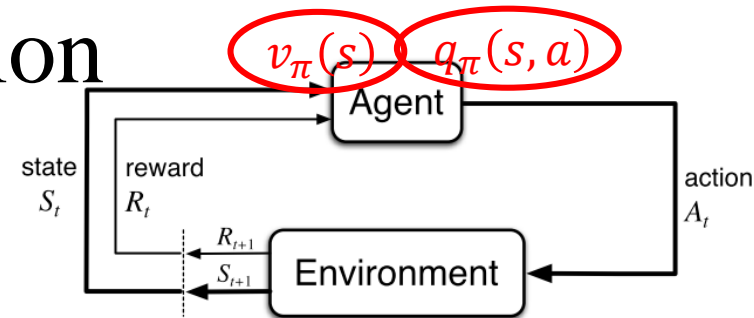
- **Value function**: how good is each state and/or action
- **Policy**: agent's behavior function
- **Model**: agent's representation of the environment

Policy

- A policy is the agent's behavior
 - It is a map from state to action,
- Policy types:
 - Deterministic policy: $a = \pi(s_i)$
 - Stochastic policy: $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$
 - ▶ Sometimes, written in $\pi(s, a)$.
- Examples:
 - In 2048: Up/down/left/right
 - In robotics: angle/force/...



Value Function




- A value function is **a prediction of future reward**
 - Used to evaluate the goodness/badness of states
 - ▶ therefore to select between actions.
 - Return $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$
- Types of value functions under policy π :
 - **State value function**: the expected return from s .

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$

$$= \mathbb{E}_\pi[G_t \mid S_t = s]$$
 - **Q-Value function**: the expected return from s taking action a .

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t \mid S_t = s, A_t = a]$$
- Examples:
 - In 2048, the expected score from a board S_t .

Model



- A **model** predicts

what the environment will do next

- \mathcal{P} is a state transition probability matrix,

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$

- ▶ predicts the next state

- \mathcal{R} is a reward function,

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$

- ▶ predicts the next (immediate) reward

- Examples:

- In 2048:

- ▶ After a move, \mathcal{P} is to generate a tile randomly as follows:

- 2-tile: with probability of 9/10

- 4-tile: with probability of 1/10



Categorizing RL Agents (Policy & Value)

- Value Based
 - No Policy (Implicit)
 - Value Function
- Policy Based
 - Policy
 - No Value Function (Implicit)
- Actor Critic
 - Policy
 - Value Function

Categorizing RL Agents (Model)

- Model Free
 - Policy and/or Value Function
 - No Model
- Model Based
 - Policy and/or Value Function
 - Model

Model-free Reinforcement Learning

- Temporal Difference (TD) Learning

- TD methods learn directly from episodes of experience
- TD is model-free: no knowledge of MDP transitions / rewards
- TD learns from incomplete episodes, by bootstrapping
- TD updates a guess towards a guess


- Monte-Carlo (MC) Learning

- MC methods learn directly from episodes of experience
- MC is model-free: no knowledge of MDP transitions / rewards
- MC learns from complete episodes: no bootstrapping
- MC uses the simplest possible idea: value = mean return
- Caveat: can only apply MC to episodic MDPs
 - ▶ All episodes must terminate
- Monte-Carlo Tree Search (MCTS) is a successful one based on MC learning.



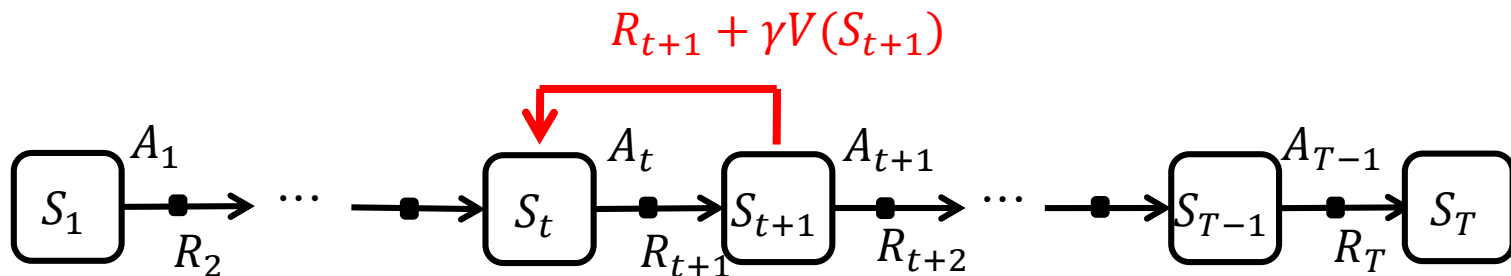
Monte-Carlo Learning

- Incremental Monte-Carlo
 - Update value $V(S_t)$ toward actual return G_t
$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$
 - α : learning rate, or called step size.
- Unbiased, but high variance.



Temporal-Difference Learning

- Simplest temporal-difference learning algorithm: TD(0)
 - Update value $V(S_t)$ toward estimated return $R_{t+1} + \gamma V(S_{t+1})$
 $V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$
 - TD target: $R_{t+1} + \gamma V(S_{t+1})$
 - TD error: $R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$
 - α : learning rate, or called step size.
- Biased, but lower variance



Case Studies

I-Chen Wu

- David Silver, Online Course for Deep Reinforcement Learning.
 - <http://www.cs.ucl.ac.uk/staff/D.Silver/web/Teaching.html>
- M. Szubert and W. Jaśkowski, “Temporal difference learning of n-tuple networks for the game 2048,” *2014 IEEE Conference on Computational Intelligence and Games (CIG)*, Aug. 2014, pp. 1–8.
- Kun-Hao Yeh, et al., Multi-Stage Temporal Difference Learning for 2048-like Games, accepted by *IEEE Transactions on Computational Intelligence and AI in Games (SCI)*, doi: 10.1109/TCIAIG.2016.2593710, 2016.
- Mnih, V. et al. Human-level control through deep reinforcement learning. *Nature* 518, 529–533 (2015).




Cases

- 2048
 - Temporal Difference (TD) Learning
 - N-tuple networks
- Atari games
 - Temporal Difference (TD) Learning
 - Deep Q-networks (DQN), a kind of Deep NN
- Go Programs (with Monte-Carlo Tree Search)
 - Monte-Carlo (MC) Learning
 - Multi-Armed Bandits
 - Planning
- AlphaGo (with Reinforcement Learning) – to be added.
 - Monte-Carlo (MC) Learning
 - Policy Gradient
- Pole Balancing – to be added.
 - Policy Gradient
 - Actor-Critic



Case Study: 2048

- [Szubert et al., 2014; Yeh et al., 2016]



2048 RL Agent

- Value function:
 - The expected score/return G_t from a board S
 - But, #states is huge
 - ▶ About 17^{16} ($=10^{20}$).
 - Empty, 2 ($=2^1$), 4 ($=2^2$), 8 ($=2^3$), ..., 65536 ($=2^{16}$).
 - Need to use value function approximator.
- Policy:
 - Simply choose the action (move) with the maximal value based on the approximator.
- Model: agent's representation of the environment
 - After a move, randomly generate a tile:
 - ▶ 2-tile: with probability of 9/10
 - ▶ 4-tile: with probability of 1/10
 - Reward: simply follow the rule of 2048.



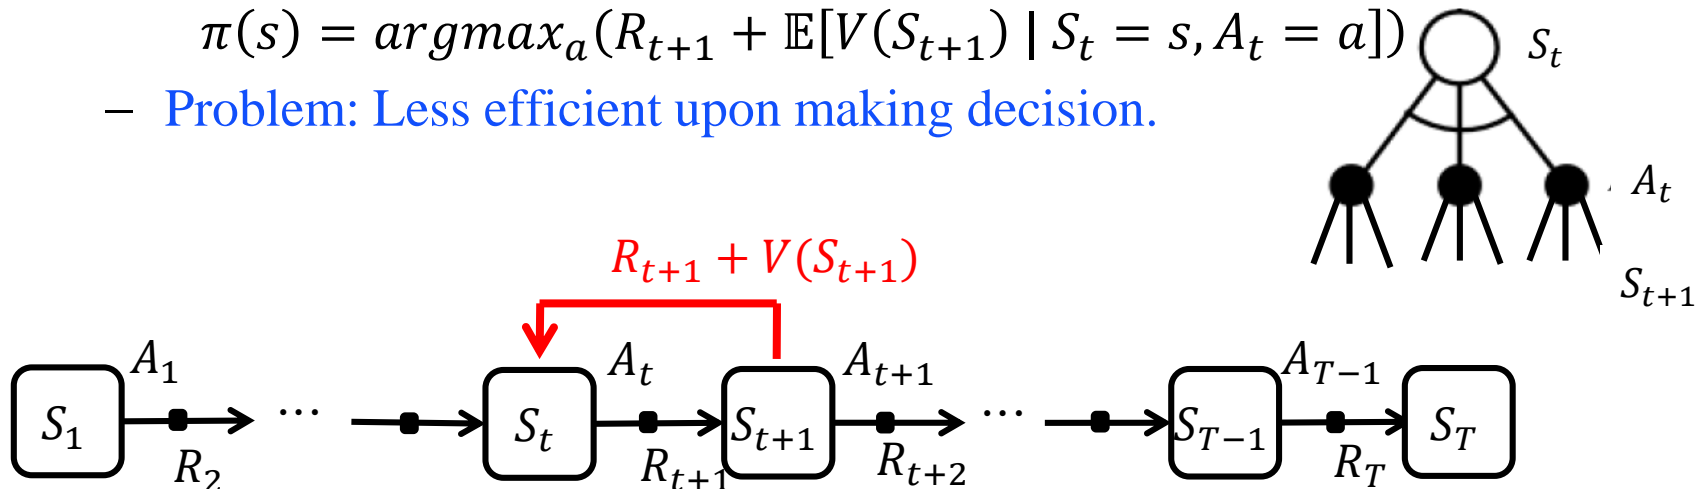
TD Learning in 2048

- State value function: (Normally $\gamma = 1$)
 - Update value $V(S_t)$ toward TD target $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$
- Making a decision (based on value).

$$\pi(s) = \operatorname{argmax}_a (R_{t+1} + \mathbb{E}[V(S_{t+1}) \mid S_t = s, A_t = a])$$

- Problem: Less efficient upon making decision.



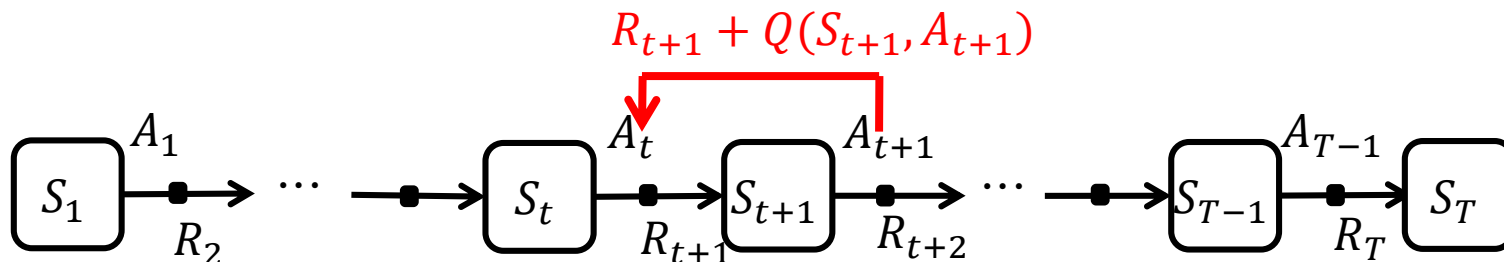

Q-Learning in 2048

- Q-value function: (Normally $\gamma = 1$)
 - Update value $Q(S_t, A_t)$ toward TD target $R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$
 - $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$

- Making decision (based on value).



$$\pi(s) = \operatorname{argmax}_a (Q(S_t, a))$$

- more efficient.
- A minor problem: Four times more memory



Afterstates in 2048

- Afterstate S_t^{af} is a state after action A_t at S_t .
 - Why not use S_t^{af} instead of (S_t, A_t) ?
 - Note: in 2048, the reward R_{t+1} is not included in S_t^{af} .
- Afterstate value function: (Normally $\gamma = 1$)
 - Update value $V^{af}(S_t^{af})$ toward TD target $\gamma \max_a (R_{t+1} + V^{af}(S_{t+1}^{af}))$
$$V^{af}(S_t^{af}) \leftarrow V^{af}(S_t^{af}) + \alpha (\gamma \max_a (R_{t+1} + V^{af}(S_{t+1}^{af})) - V^{af}(S_t^{af}))$$
- Making decision (based on value).
 - $\pi(s) = \operatorname{argmax}_a (V^{af}(S_t^{af}))$
 - For simplicity, we use V , instead of V^{af} , if it can be applied to both.
 - More efficient.



Value Function Approximation

- As mentioned above, #states is huge, so we need to use value function approximation.
 - Use a value function approximator, $\hat{v}(S, \theta) \approx V(S)$.
 - Simply use **deterministic policy**: $\pi(S) = \operatorname{argmax}_a(\hat{v}(S, \theta))$
- But, what kind of value function approximator can we use?
 - What features can we choose?
 - ▶ Traditionally, # of empty cells, # of continuous cells, big tiles, etc.
 - **Linear** (like n-tuple network) vs. **non-linear** (like NN)
- n-tuple network is a powerful network for 2048.
 - Explore **a large set of features**.
 - Simplify operations by **linear value function approximation**.



Linear Value Function Approximation

- Represent value function by a linear combination of features

$$\hat{v}(S; \theta) = x(S)^T \theta = \sum_{j=1}^n x_j(S) \theta_j$$

- Gradient of $\hat{v}(S, \theta)$:

$$\nabla_{\theta} \hat{v}(S, \theta) = x(S)$$



Gradient Descent

- Update value $V(S_t)$ towards TD target $y_t = R_{t+1} + V(S_{t+1})$

$$\Delta V = (R_{t+1} + V(S_{t+1}) - V(S_t)) = (y_t - V(S_t))$$

$$V(S_t) \leftarrow V(S_t) + \alpha \Delta V$$

- α : learning rate, or called step size.
- Note: $\gamma = 1$ here.

- Objective function is to minimize the following loss in parameter θ . (note: $\hat{v}(S, \theta) = x(S)^T \theta$)

$$\mathcal{L}(w) = \mathbb{E} \left[(y_t - \hat{v}(S, \theta))^2 \right]$$

$$\nabla_{\theta} \mathcal{L}(\theta) = (y_t - \hat{v}(S, \theta)) \cdot \nabla_{\theta} \hat{v}(S, \theta) = \Delta V \cdot x(S)$$

- Update features w : step-size * prediction error * feature value
 $\theta \leftarrow \theta + \alpha \Delta V \cdot x(S)$



N-Tuple Network

- Example: 4-tuple networks as shown.
 - Each cell has 16 different tiles
 - 16^4 features for this network.
 - ▶ But only one is on, others are 0.
 - ▶ So, we can use table lookup to find the feature weight.



64	⁰	8	4
128	2 ¹		2
2	8 ²		2
128	³		

0123	weight
0000	3.04
0001	-3.90
0002	-2.14
⋮	⋮
0010	5.89
⋮	⋮
0130	-2.01
⋮	⋮



Other N-Tuple Networks

- Left: [Szubert et al., 2014]; Right: [Yeh et al., 2016]
- Some researchers even used 7-tuple network.



Update Features in N-Tuple Networks

- For n-tuple networks, simply update values with $\alpha\Delta V$ at $LUT_i[index(s_i)]$
- Features:
 - 8 x 16^4 features, $x(S) = [0, 1, 0, \dots, 0, 0, 1, \dots, \dots, 1, 0, 0, \dots]$
 - ▶ All 0s, except for 8 ones.
 - One 1 every 16^4 features.
 - Let their indices be $s_1, s_2, s_3, s_4, s_5, s_6, s_7, s_8$.
 - Only need to update $\alpha\Delta V$ at the features indexed by these indices.
 - Very efficient and fast.
- For k n-tuple networks,

$$\hat{v}(S, \theta) = x(S)^T \theta = \sum_{j=1}^n x_j(S) \theta_j = \sum_{i=1}^k LUT_i[index(s_i)]$$
 - LUT_i : the i-th n-tuple network lookup table.
 - $index(s_i)$: The index in the i-th n-tuple network of state S .
- Update features w : step-size * prediction error * feature value
 - $\theta \leftarrow \theta + \alpha\Delta V \cdot x(S)$
 - Only need to update values θ_j with $\alpha\Delta V$ at $LUT_i[index(s_i)]$.



Afterstate Evaluation Function

```
1: function EVALUATE( $s, a$ )
2:    $s', r \leftarrow \text{COMPUTE AFTERSTATE}(s, a)$ 
3:   return  $r + V(s')$ 
4:
5: function LEARN EVALUATION( $s, a, r, s', s''$ )
6:    $a_{next} \leftarrow \arg \max_{a' \in A(s'')} \text{EVALUATE}(s'', a')$ 
7:    $s'_{next}, r_{next} \leftarrow \text{COMPUTE AFTERSTATE}(s'', a_{next})$ 
8:    $V(s') \leftarrow V(s') + \alpha(r_{next} + V(s'_{next}) - V(s'))$ 
```




```
1: function PLAY GAME
2:    $score \leftarrow 0$ 
3:    $s \leftarrow \text{INITIALIZE GAME STATE}$ 
4:   while  $\neg \text{IS TERMINAL STATE}(s)$  do
5:      $a \leftarrow \arg \max_{a' \in A(s)} \text{EVALUATE}(s, a')$ 
6:      $r, s', s'' \leftarrow \text{MAKE MOVE}(s, a)$ 
7:     if LEARNING ENABLED then
8:       LEARN EVALUATION( $s, a, r, s', s''$ )
9:      $score \leftarrow score + r$ 
10:     $s \leftarrow s''$ 
11:   return  $score$ 
12:
13: function MAKE MOVE( $s, a$ )
14:    $s', r \leftarrow \text{COMPUTE AFTERSTATE}(s, a)$ 
15:    $s'' \leftarrow \text{ADD RANDOM TILE}(s')$ 
16:   return ( $r, s', s''$ )
```





The N-Tuple Networks Used

- Use the following [Szubert and Jaskowski 2014]



- Ours:



Performance Results (without search)

2048 rate	100%
4096 rate	100%
8192 rate	99.20%
16384 rate	83.30%
32768 rate	8.10%
Maximum score	607488
Average score	331820



Case Study: Atari 2600 Games

- Learn to play Atari games **from video only** (without knowing the game a priori)



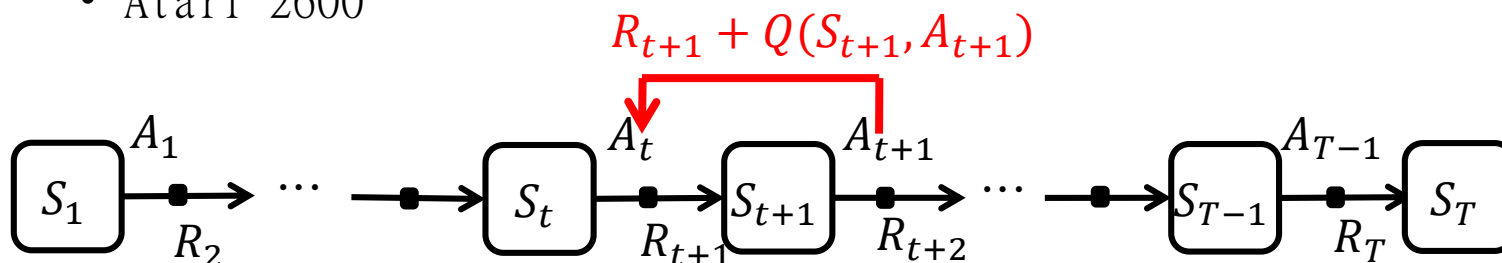
- Atari 2600



- Breakout



- Space Invaders



Deep Q-Networks (DQN)



DQN uses experience replay and fixed Q-targets

- Take action according to ϵ -greedy policy
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory \mathcal{D}
- Sample random mini-batch of transitions (s, a, r, s') from \mathcal{D}
- Compute Q-learning targets w.r.t. old, fixed parameters θ^-
- Optimize MSE between Q-network and Q-learning targets
 - Minimize a sequence of loss functions $\mathcal{L}(\theta_i)$ that changes at each iteration i .
 - $\mathcal{L}_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}_i} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$
- Using variant of stochastic gradient descent
 - Differentiating the loss function with respect to the weights we arrive at the following gradient
 - $\nabla_{\theta_i} \mathcal{L}_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}_i} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right) \cdot \nabla_{\theta_i} Q(s, a; \theta_i) \right]$




DQN in Atari

- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s
 - stack of raw pixels from last 4 frames
- Output
 - $Q(s, a_i | \theta)$ for 18 joystick/button positions
- Reward
 - change in score for that step




Performance of Deep Q-Learning

- Left (**stronger than human**)




Case Study: Go (MCTS)

- Monte-Carlo Tree Search:
 - Monte-Carlo (MC) Learning (**z: 1 for win, 0 for loss**)
 - Multi-Armed Bandits
 - **Planning**
- Very successful for Go in the past decade.
- And also applied to others successfully.
 - Other games like Havannah, Hex, GGP
 - Other applications, like mathematical optimization problems (scheduling, UCP, camera coverage).




Go – One of the Most Popular Games

- Game tree complexity: about 10^{360}
 - It is just impossible to try all moves.



Can Alpha-Beta Search Work for Go?

- Alpha-Beta Search
 - Very successful for many games such as **chess**.
 - ▶ **Almost dominate all computer games before 2006.**
 - ▶ This is what Deep Blue used.
- The key for chess: evaluate position accurately and efficiently.
E.g., features:
 - King: 1000
 - Queen: 200
 - Rook: 100
 - Knight: 80
 - Bishop: 70
 - Pawn: 30
 - Guarded Pawns: 30
 - Guarded Knights: 40
 - ...
- Problem for chess:
 - need to **consult with experts for feature values.**




Why not alpha-beta search for Go?

- No simple heuristics like chess.
 - Only black/white pieces (no difference)
- Must know life-and-death
 - But, all are correlated.
 - ▶ Like the lower-right one.
 - But, this is very complex.


Since no simply heuristics to evaluate,

- Why not use Monte-Carlo?
- Calculate it based on stochastics.



Rules Overview Through a Game (opening 1)

- Black/White move alternately by putting one stone on an intersection of the board.




The example was given by B. Bouzy at CIG'07.



Rules Overview Through a Game

(opening 2)


- Black and White aims at surrounding large « zones »



Rules Overview Through a Game


(atari 1)

- A white stone is put into « atari » : it has only one liberty left.




Rules Overview Through a Game (defense)

- White plays to connect the one-liberty stone yielding a four-stone white string with 5 liberties.




Rules Overview Through a Game (atari 2)

- It is White's turn. One black stone is atari.



Rules Overview Through a Game


- White plays on the last liberty of the black stone which is removed



Rules Overview Through a Game


(human end of game)

- The game ends when the two players pass. (Experts would stop here)




Rules Overview Through a Game (contestation 1)

- White contests the black « territory » by playing inside.



Rules Overview Through a Game (contestation 2)


- White contests black territory, but the 3-stone white string has one liberty left



Rules Overview Through a Game

(follow up 1)


- Black has captured the 3-stone white string



Rules Overview Through a Game

(follow up 2)


- White lacks liberties...



Rules Overview Through a Game

(follow up 3)


- Black suppresses the last liberty of the 9-stone string
- Consequently, the white string is removed



Rules Overview Through a Game

(follow up 4)


- Contestation is going on. White has captured four black stones.



Rules Overview Through a Game



(concrete end of game)

- The board is covered with either stones or « eyes ». Programs know to end.



Stochastics

- Calculate values based on stochastics.
 - Good example: calculate π .



Multi-Armed Bandit Problem

(吃角子老虎問題)

- Assume that you have infinite plays
 - How to choose the one with the maximal average return?



Exploration vs. Exploitation

● Example for the exploration vs exploitation dilemma

- **Exploration:** is a long-term process, with a risky, uncertain outcome.
- **Exploitation:** by contrast is short-term, with immediate, relatively certain benefits



Deterministic Policy: UCB1

- UCB: Upper Confidence Bounds. [Auer *et al.*, 2002]
- Observed rewards when playing machine i : $X_{i,1}, X_{i,2}, \dots$
- Initialization: Play each machine once.

- Loop:
 - Play machine j that maximizes, $\bar{X}_j + \sqrt{\frac{2 \log n}{T_j(n)}}$

where n is the overall number of plays done so far,


$$\bar{X}_{i,s} = \frac{1}{s} \sum_{j=1}^s X_{i,j} \quad , \quad \bar{X}_i = \bar{X}_{i,T_i(n)} \quad ,$$

- Key:
 - Ensure optimal machine is played exponentially more often than any other machine.




Cumulative Regret

- Assume Machines M_1, M_2, M_3, M_4, M_5
 - Win rates: 37%, 42%, 47%, 52%, 57%
 - Trial numbers: n_1, n_2, n_3, n_4, n_5 .




Monte-Carlo Tree Search

- A kind of planning
- A kind of **Reinforcement learning**



Strength of Go Program after MCTS

- [Schaeffer et al., 2014]



Strength grew fast, after MCTS.




Case Study: AlphaGo

- Use **stochastic policy gradient ascent** to maximize the likelihood of the human move a selected in state s

$$\Delta\theta = \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) \cdot z$$

- θ : network parameter.
- α : learning rate
- z : the value of the episode
 - ▶ win/loss (1/-1) of the game



Case Study: Pole Balancing

[Lillicrap 2016] Continuous control with deep reinforcement learning, ICLR, 2016 (DeepMind),

- Deep Policy Gradient

- Actor-critic, model-free algorithm
- underlying the success of Deep Q-Learning to the continuous action domain

$$\Delta\theta = \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) \cdot Q_w(s_t, a_t)$$

